

NYU FRE 7773 - Week 7

Machine Learning in Financial Engineering


Ethan Rosenthal

Time Series Machine Learning

Machine Learning in Financial Engineering


Ethan Rosenthal

Time



Uber trips


<https://eng.uber.com/neural-networks/>



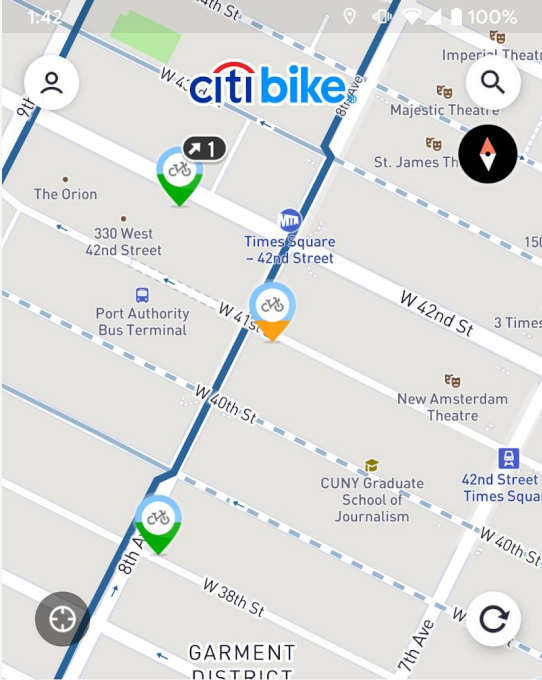
Amazon weekly item sales

DeepAR: Probabilistic Forecasting with Autoregressive Recurrent Networks

<https://arxiv.org/abs/1704.04110>



No Citi Bikes



0 seconds ago

W 41 St & 8 Ave

2

Bikes

55

Docks

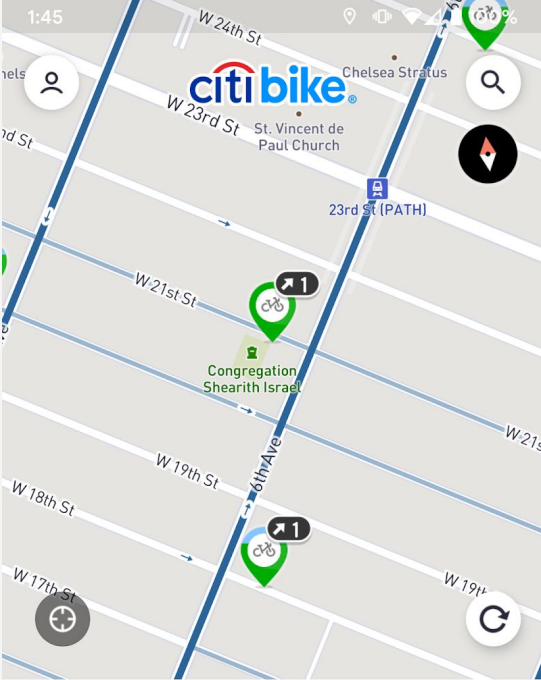


Unlock a Bike



Plan a Ride

No Docks



13 seconds ago

W 21 St & 6 Ave

49

Bikes

0


Docks




Unlock a Bike



Plan a Ride






“Classical” Time Series Modeling



State Space Models



Nonlinear PDE's



ARIMA



Gauss-Markov Theorem




$$\begin{array}{ccccc}
 (1 - \phi_1 B - \dots - \phi_p B^p) & (1 - B)^d y_t & = & c + (1 + \theta_1 B + \dots + \theta_q B^q) \varepsilon_t \\
 \uparrow & \uparrow & & \uparrow \\
 \text{AR}(p) & d \text{ differences} & & \text{MA}(q)
 \end{array}$$




**Can you skip
all of this?**

Inference



Prediction

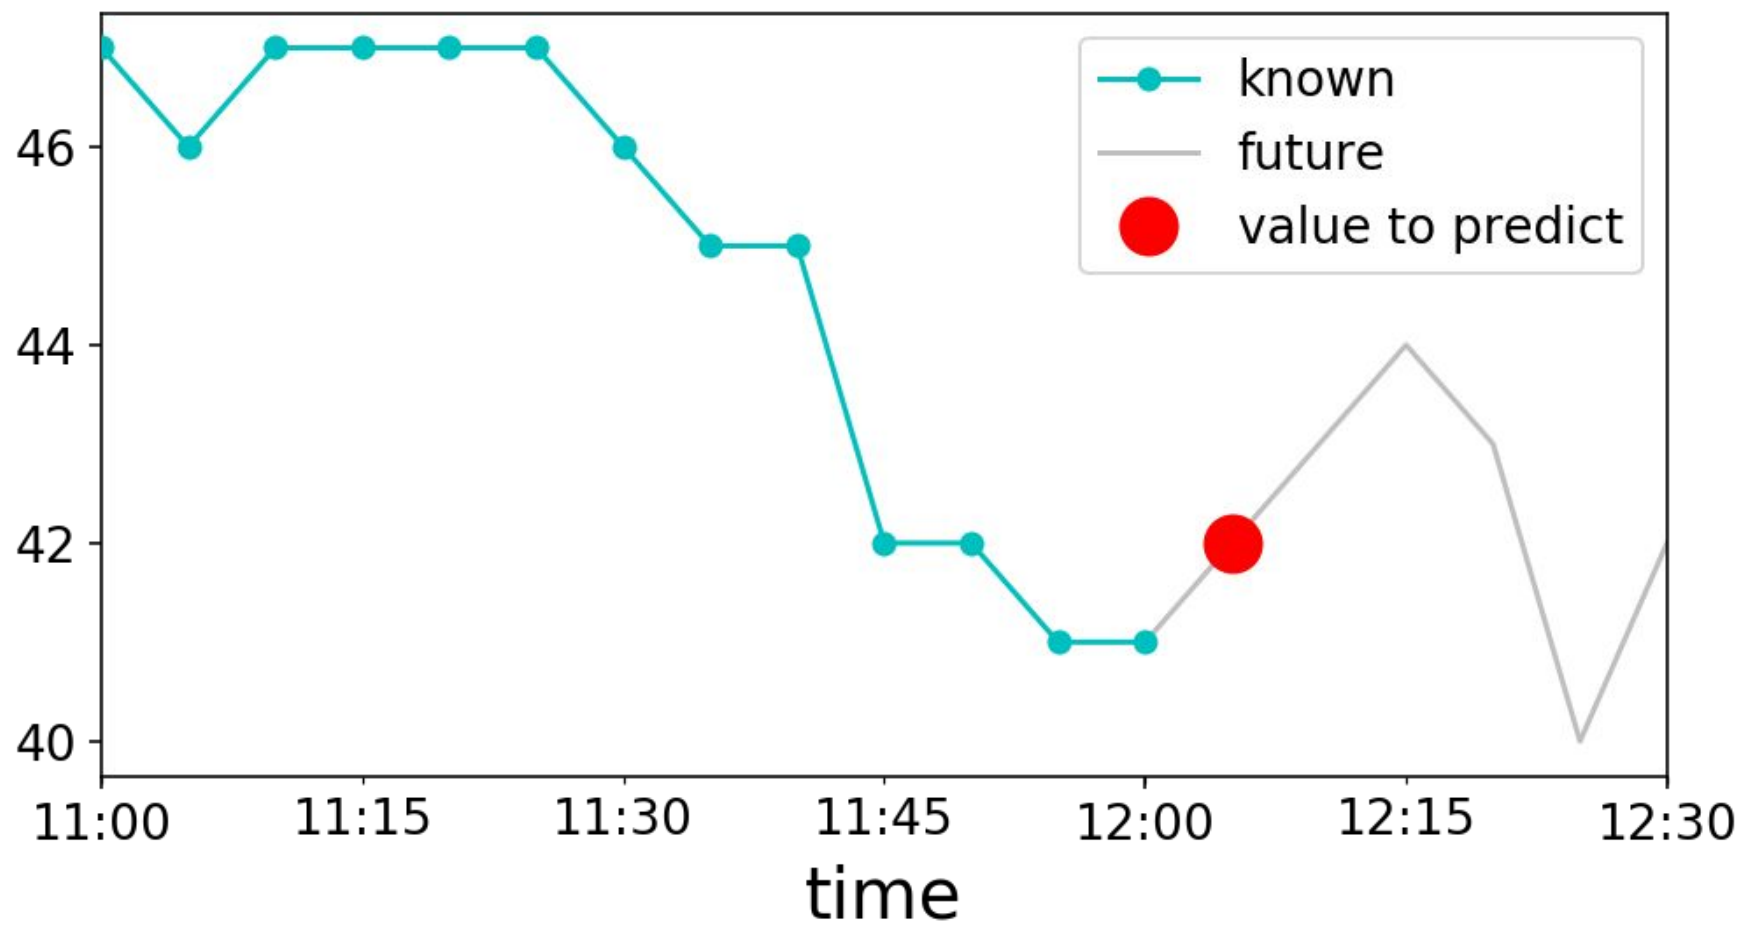


Where's the X Matrix?

```
f.fit(X, y)  
y_new = f.predict(X_new)
```

In the beginning, there was **y**

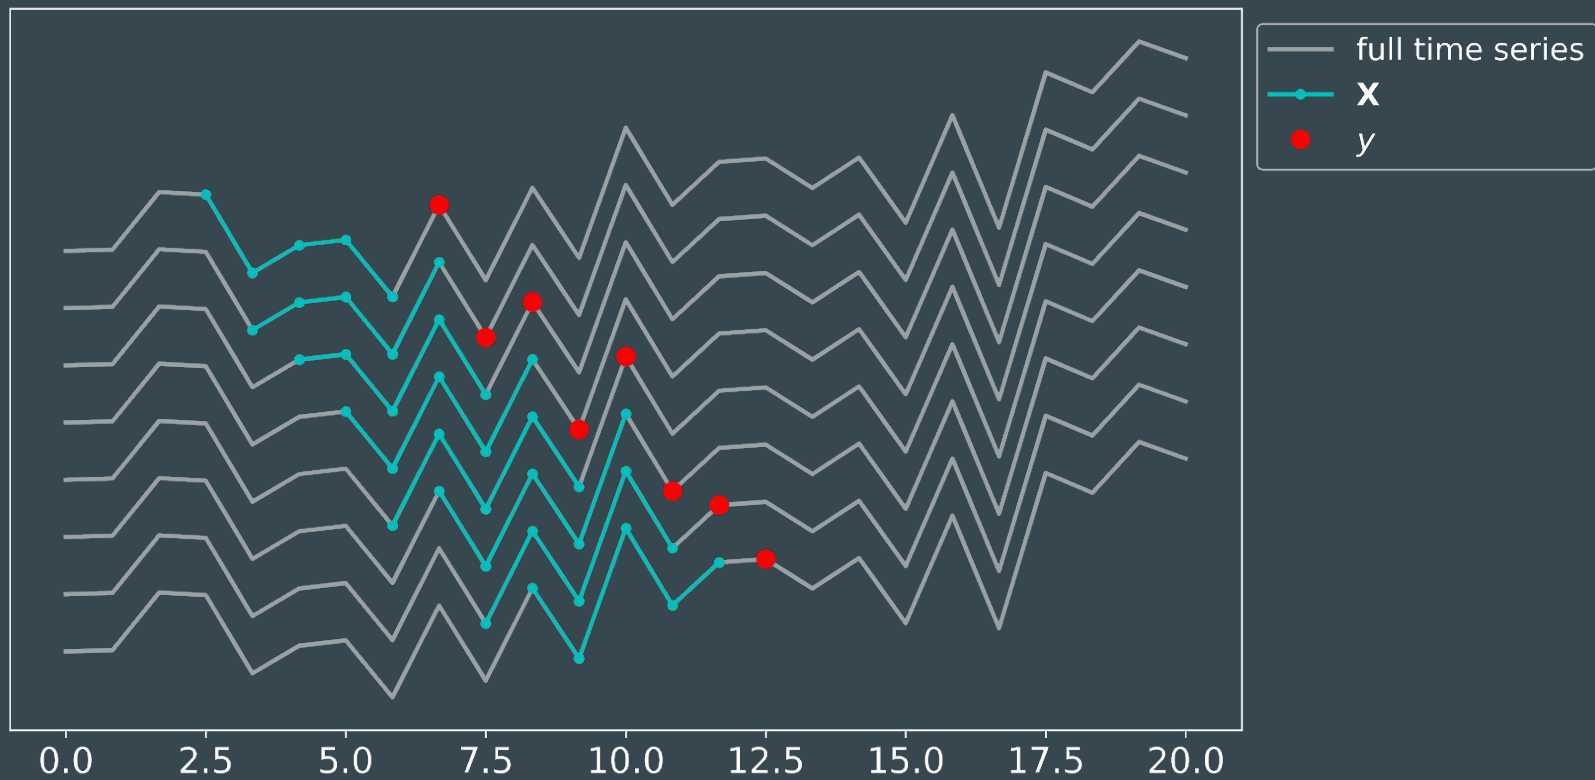




```
X = []  
for idx in range(len(y) - window):  
    X.append(y[idx:idx + window])  
  
X = np.array(X)
```

```
y: array([0,  
         1,  
         2,  
         3,  
         4,  
         5,  
         6,  
         7,  
         8,  
         9,  
        10])
```

```
X: array([[0, 1, 2, 3, 4],  
         [1, 2, 3, 4, 5],  
         [2, 3, 4, 5, 6],  
         [3, 4, 5, 6, 7],  
         [4, 5, 6, 7, 8],  
         [5, 6, 7, 8, 9]])
```



$$\mathbf{X}\beta = \hat{\mathbf{y}}$$

$$\begin{bmatrix} y_0 & y_1 & y_2 & \cdots & y_{w-1} \\ y_1 & y_2 & y_3 & \cdots & y_w \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ y_{t-2-w} & y_{t-1-w} & y_{t-w} & \cdots & y_{t-2} \\ y_{t-1-w} & y_{t-w} & y_{t-w+1} & \cdots & y_{t-1} \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_{w-2} \\ \beta_{w-1} \end{bmatrix} = \begin{bmatrix} \hat{y}_w \\ \hat{y}_{w+1} \\ \vdots \\ \hat{y}_{t-1} \\ \hat{y}_t \end{bmatrix}$$

Modeling

- Not limited to Linear Regression. Use trees, neural nets, whatever you want!
- Not limited to regression. Classification, quantile regression, etc...

Preprocessing and Feature Engineering

- Differencing
- Rolling Mean
- Filters
- Fourier components
- Seasonal lags
- etc...

Adding “Exogenous” Features

```
X = []  
for idx in range(len(y) - window):  
    X.append(y[idx:idx + window])
```

```
X = np.array(X)  
X = np.hstack((X, X_features))
```

X: array([[0, 1, 2, 3, 4, 0.5, -0.1],
[1, 2, 3, 4, 5, 2.3, 0.2],
[2, 3, 4, 5, 6, -0.2, 0.4],
[3, 4, 5, 6, 7, 0.9, 1.1],
[4, 5, 6, 7, 8, 1.2, 0.5],
[5, 6, 7, 8, 9, -0.7, -0.2]])

Lag Features **Extra Features**

y: array([0,
1,
2,
3,
4,
5,
6,
7,
8,
9,
10])

```
X = []  
for idx in range(len(y) - window):  
    X.append(y[idx:idx + window])
```

```
X = np.array(X)  
X = np.hstack((X, X_features))
```

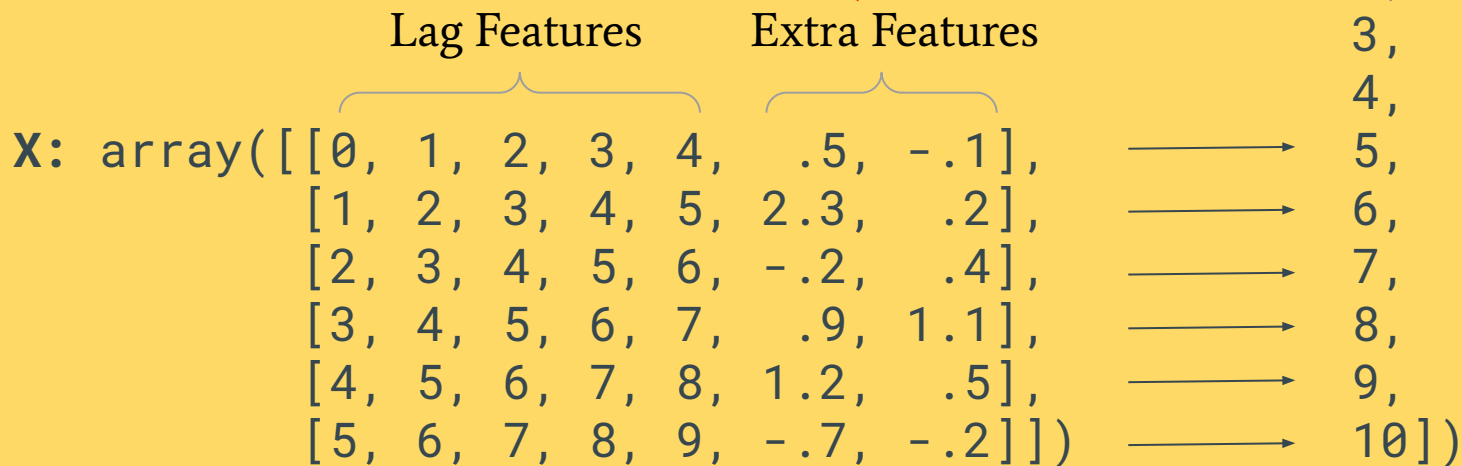
Adding “Exogenous” Features

Complicated to
construct!

Lag Features **Extra Features**

X: array([[0, 1, 2, 3, 4, .5, -.1],
[1, 2, 3, 4, 5, 2.3, .2],
[2, 3, 4, 5, 6, -.2, .4],
[3, 4, 5, 6, 7, .9, 1.1],
[4, 5, 6, 7, 8, 1.2, .5],
[5, 6, 7, 8, 9, -.7, -.2]])

y: array([0,
1,
2,
3,
4,
5,
6,
7,
8,
9,
10])



Recap

- Take your time series \mathbf{y} .
- Treat each point in \mathbf{y} as a point that you want to predict.
- Construct \mathbf{X} from any data you want that comes *prior* to the point in \mathbf{y} that you want to predict.

$$\mathbf{X}_t = \mathbf{y}_{t' < t}$$

- Fit a regression model on \mathbf{X} and \mathbf{y} .
- For each point in \mathbf{y} , use model to predict the next point.


$$\hat{y}_t = f(\mathbf{X}_t)$$

$$\hat{y}_t = f(\mathbf{y}_{t' < t-1})$$


Forecasting




Recursive Forecasting




Recursive Forecasting



Recursive Forecasting



Recursive Forecasting




```
def recursive_forecast(model, input_data, num_points_in_future):  
    for point in range(num_points_in_future):  
        prediction = model.predict(input_data)  
        # Append prediction to the input data  
        input_data = np.hstack((input_data, prediction))  
  
    return prediction
```

Optimize for
next step


Pray recursive
steps work

Horizon Forecasting



More info: “Machine learning strategies for multi-step-ahead time series forecasting”,
Souhaib Ben Taieb https://souhaib-bentaieb.com/papers/2014_phd.pdf

Horizon Forecasting



```
y: array([ 0.,  
          1.,  
          2.,  
          3.,  
          4.,  
          5.,  
          6.,  
          7.,  
          8.,  
          9.,  
         10.])  
y_horizon: array([[ 0.,  1.,  2.],  
                  [ 1.,  2.,  3.],  
                  [ 2.,  3.,  4.],  
                  [ 3.,  4.,  5.],  
                  [ 4.,  5.,  6.],  
                  [ 5.,  6.,  7.],  
                  [ 6.,  7.,  8.],  
                  [ 7.,  8.,  9.],  
                  [ 8.,  9., 10.],  
                  [nan, nan, nan],  
                  [nan, nan, nan]])
```


Predicting multiple targets


- [sklearn.multioutput.MultiOutputRegressor](#)
- Train an individual model for each target.
- Pros:
 - Very simple
 - Works with any model
- Cons
 - Resource intensive
 - No sharing of knowledge

Predicting multiple targets

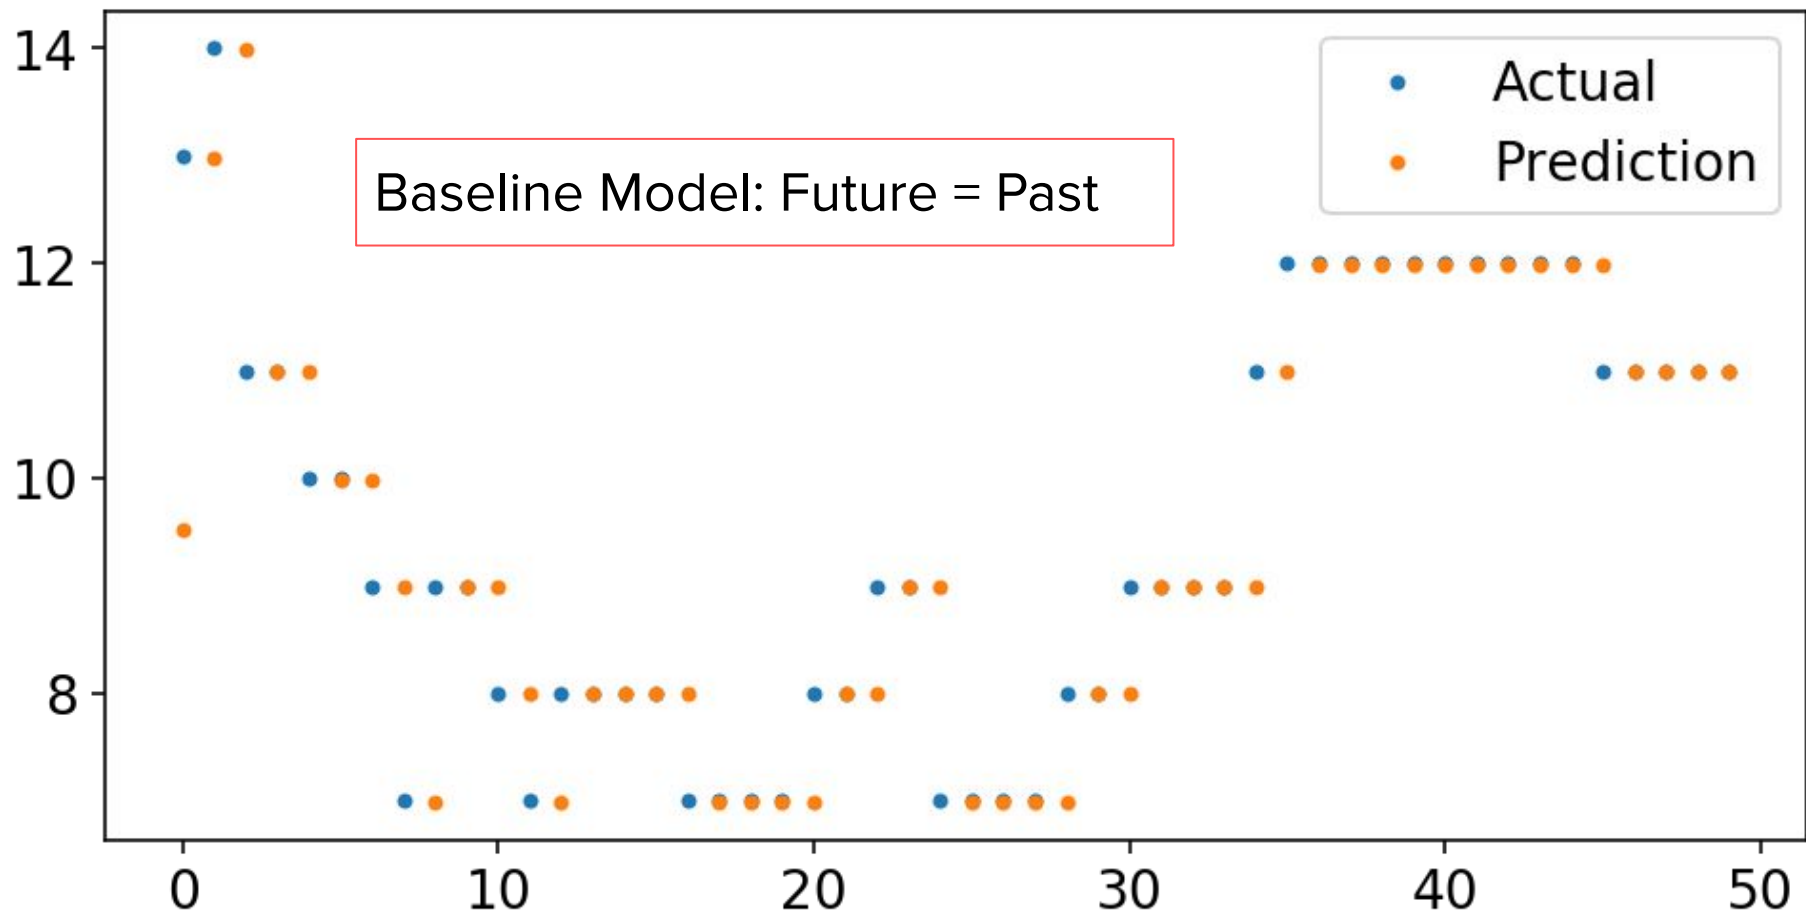
- Deep learning model with multiple outputs.
- Pros:
 - Potentially less resource intensive
 - Direct optimization
 - Sharing of knowledge
- Cons:
 - All the caveats of deep learning
 - Limited to the horizon

Evaluation







Baseline Model: Future = Past




Forecasting Views




Forecasting Views




Forecasting Views



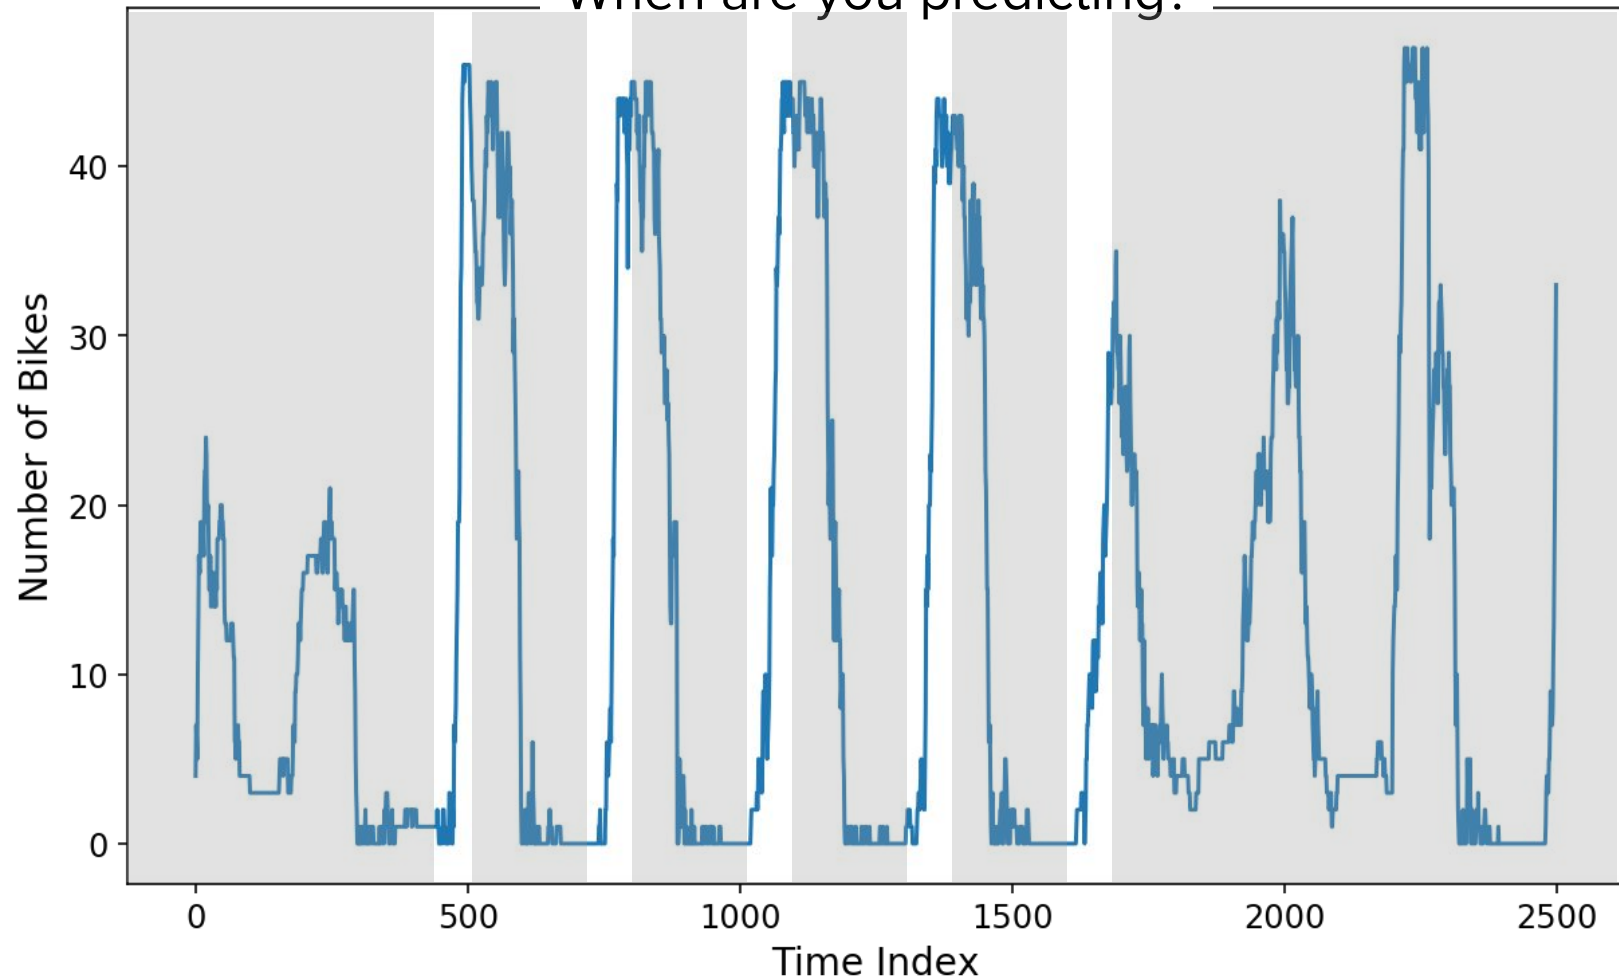
Metrics Aggregation




When are you predicting?




When are you predicting?



Regression or Classification?



What do you know when
you predict?



Number of Bikes

Scatter plot showing the relationship between the date (last_reported) and the count. The x-axis represents the date from 10-13 00 to 10-13 09. The y-axis represents the count, ranging from 0.0 to 17.5. The data points show a sharp increase in counts starting around 10-13 07, peaking at approximately 17.5 around 10-13 09.

What do you know when
you predict?

