

NYU FRE 7773 - Week 11

Machine Learning in Financial Engineering

Jacopo Tagliabue

Serving predictions

Machine Learning in Financial Engineering
Jacopo Tagliabue

Welcome to the jungle

If your work needs to have an impact, it needs to RUN OUTSIDE YOUR LAPTOP:

1. Your code can be **inspected, modified, understood** by others, typically your technical colleagues: you need to write clean, modular, testable code and make your pipeline fully reproducible.
2. Your model can be **trusted** by others, typically, other stakeholders, who may or may not be technical folks: you need to “make sure” the model behaved as designed before pushing it in front of end-users.
3. Predictions can be **consumed** by others, typically anybody with an internet connection: you need to expose your model as an endpoint which returns predictions when supplied with the appropriate parameters.

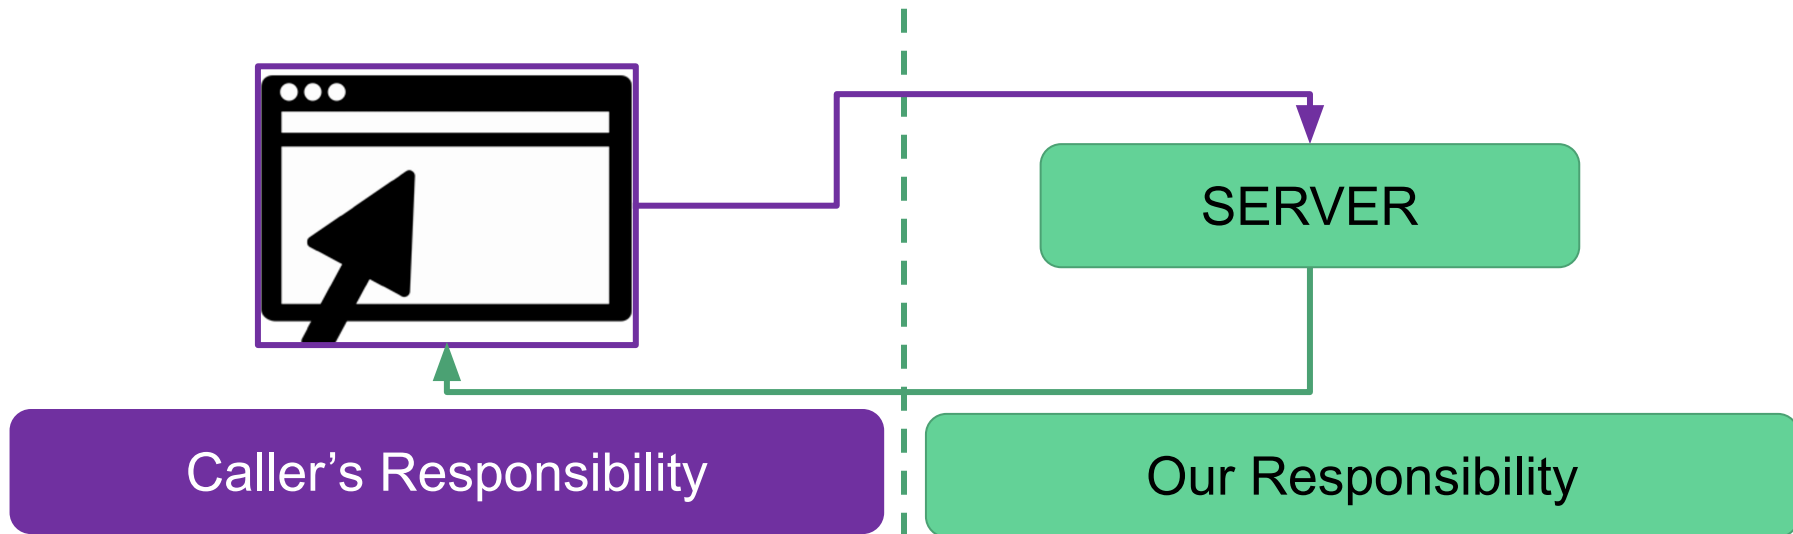
Welcome to the jungle

If your work needs to have an impact, it needs to RUN OUTSIDE YOUR LAPTOP:

1. Your code can be **inspected, modified, understood** by others, typically your technical colleagues: you need to write clean, modular, testable code and make your pipeline fully reproducible.
2. Your model can be **trusted** by others, typically, other stakeholders, who may or may not be technical folks: you need to “make sure” the model behaved as designed before pushing it in front of end-users.
3. Predictions can be **consumed** by others, typically anybody with an internet connection: you need to expose your model as an endpoint which returns predictions when supplied with the appropriate parameters.


Part 3: Serving predictions

- If our model stays on our laptop, nobody will be able to use it!
- **Client-server architecture:** our model interacts with *many* remote clients through an API (also called “endpoint”) - we abstract away model code (and complexity) and expose a pure input-output interface: clients send us the input, we return a prediction.



Part 3: Serving predictions

- If our model stays on our laptop, nobody will be able to use it!
- **Client-server architecture:** our model interacts with *many* remote clients through an [API](#) (also called “endpoint”) - we abstract away model code (and complexity) and expose a pure input-output interface: clients send us the input, we return a prediction.



HTML (+ CSS) + Javascript

SERVER

Python ([Flask](#), [FastAPI](#) etc.)

Show me first!

Intro to Flask applications

- We will be using Flask as a simple framework to serve model predictions after training
- Flask has several attractive features:
 - Helps with structuring both the front-end (the web page) and back-end (the endpoint)
 - Pure Python back-end
 - Minimal syntax for routing, GET / POST etc.

Step 1: prepare a web page

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4   <title>{{ project }} app</title>
5   <script src="https://ajax.googleapis.com/ajax/libs/jquery/3.5.1/jquery.min.js"></script>
6 </head>
7 <script type="text/javascript">
8
9   $(function() {
10     $('#predict').click(function() {
11       event.preventDefault();
12       var form_data = new FormData($('#myform')[0]);
13       console.log(form_data);
14       $.ajax({
15         type: 'POST',
16         url: '/',
17         data: form_data,
18         contentType: false,
19         processData: false,
20       }).done(function(data, textStatus, jqXHR){
21         $('#result').text(data);
22       }).fail(function(data){
23         alert('error!');
24       });
25     });
26   });
27
28 </script>
29 <body>
30   <h1>{{ project }}</h1>
```

- Prepare a simple HTML page for users to interact with our endpoint.
 - It is not much different than the streamlit app we built before!
- Note that we use a simple Javascript function with jQuery to perform a POST request.

Step 2: prepare the Flask back-end application

```
32 # We need to initialise the Flask object to run the flask app
33 # By assigning parameters as static folder name, templates folder name
34 app = Flask(__name__, static_folder='static', template_folder='templates')
```

- Initialize a Flask app in [app.py](#)
 - Note the *templates* folder contains the HTML we created before!
- Make sure the app is started when we run “flask run”: the script will spin up a web server that will be ready to listen for incoming requests (from our HTML page, of course)

```
54 if __name__=='__main__':
55     # Run the Flask app to run the server
56     app.run(debug=True)
```

Step 3: load the ML model in memory

```
13 ##### THIS IS GLOBAL, SO OBJECTS LIKE THE MODEL CAN BE RE-USED ACROSS REQUESTS #####
14
15 FLOW_NAME = 'MyRegressionFlow' # name of the target class that generated the model
16 # Set the metadata provider as the src folder in the project,
17 # which should contains /.metaflow
18 metadata('../src')
19 # Fetch currently configured metadata provider to check it's local!
20 print(get_metadata())
21
22 def get_latest_successful_run(flow_name: str):
23     "Gets the latest successfull run."
24     for r in Flow(flow_name).runs():
25         if r.successful:
26             return r
27
28 # get artifacts from latest run, using Metaflow Client API
29 latest_run = get_latest_successful_run(FLOW_NAME)
30 latest_model = latest_run.data.model
```

- As in all Python scripts, what is declared outside the functions is “global”.
 - In this context, this means that objects can live *across requests*: if the client ask for prediction 1 and then prediction 2, we do not need to reload the model, as it is already **in memory!**
- We want the model to be “global” as retrieving the model from Metaflow storage may be slow and expensive.


Step 4: defining our endpoint

```
36 @app.route('/', methods=['POST', 'GET'])
37 def main():
38
39     # on GET we display the page
40     if request.method=='GET':
41         return render_template('index.html', project=FLOW_NAME)
42     # on POST we make a prediction over the input text supplied by the user
43     if request.method=='POST':
44         # debug
45         # print(request.form.keys())
46         _x = request.form['_x']
47         val = latest_model.predict([[float(_x)]])
48         # debug
49         print(_x, val)
50         # Returning the response to the client
51         return "Predicted Y is {}".format(val)
```

- We use a decorator to define our route (empty in this case, but could be, say, “predict”).
 - If it was “predict”, our server would be listening for calls at **URL/predict**
- We distinguish between page load (GET) and request from a prediction (POST).

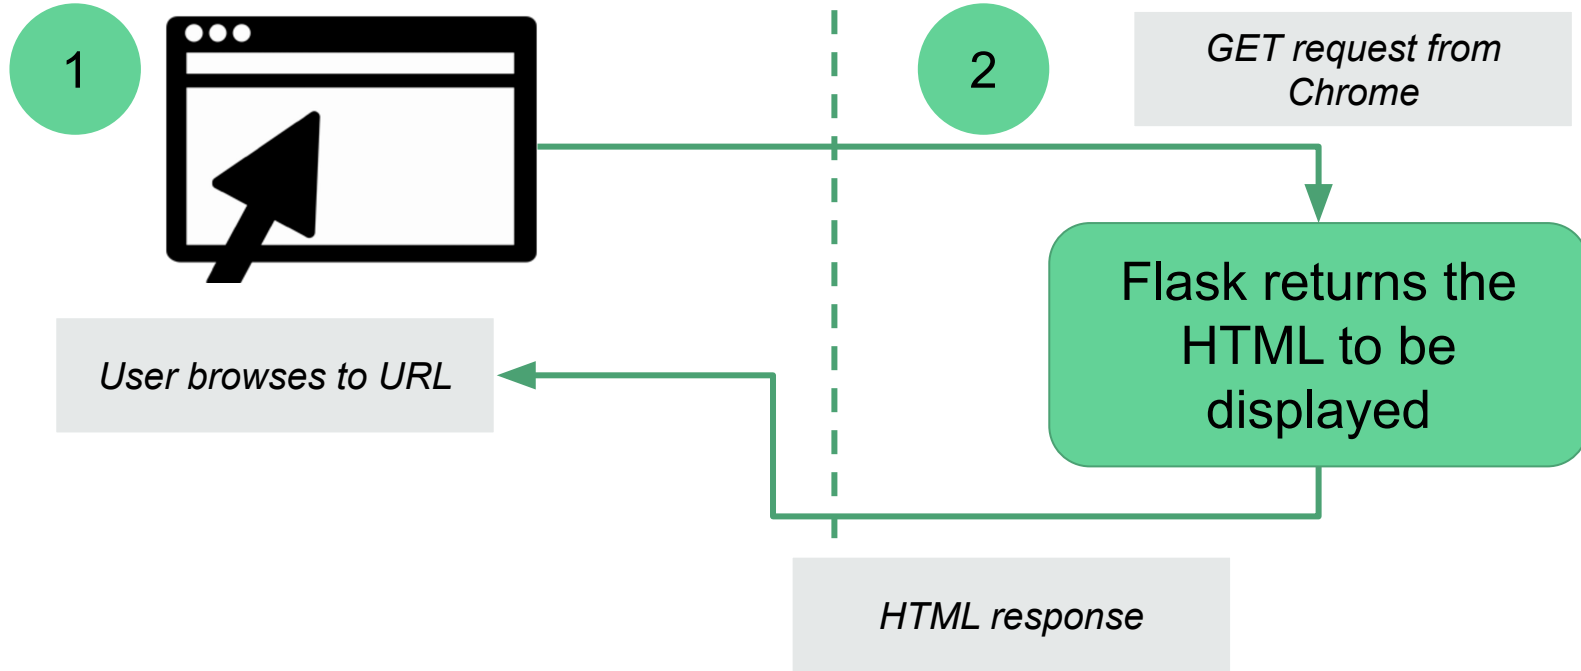
The full workflow

1




User browses to URL

The full workflow



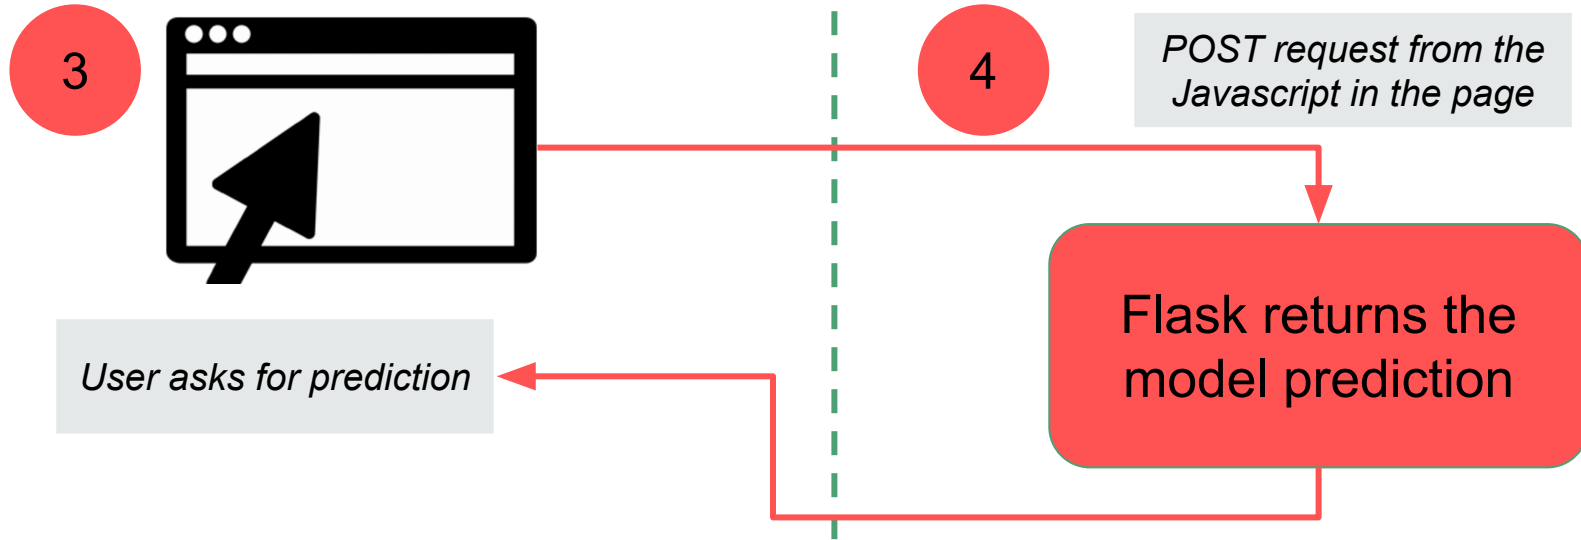
The full workflow

3



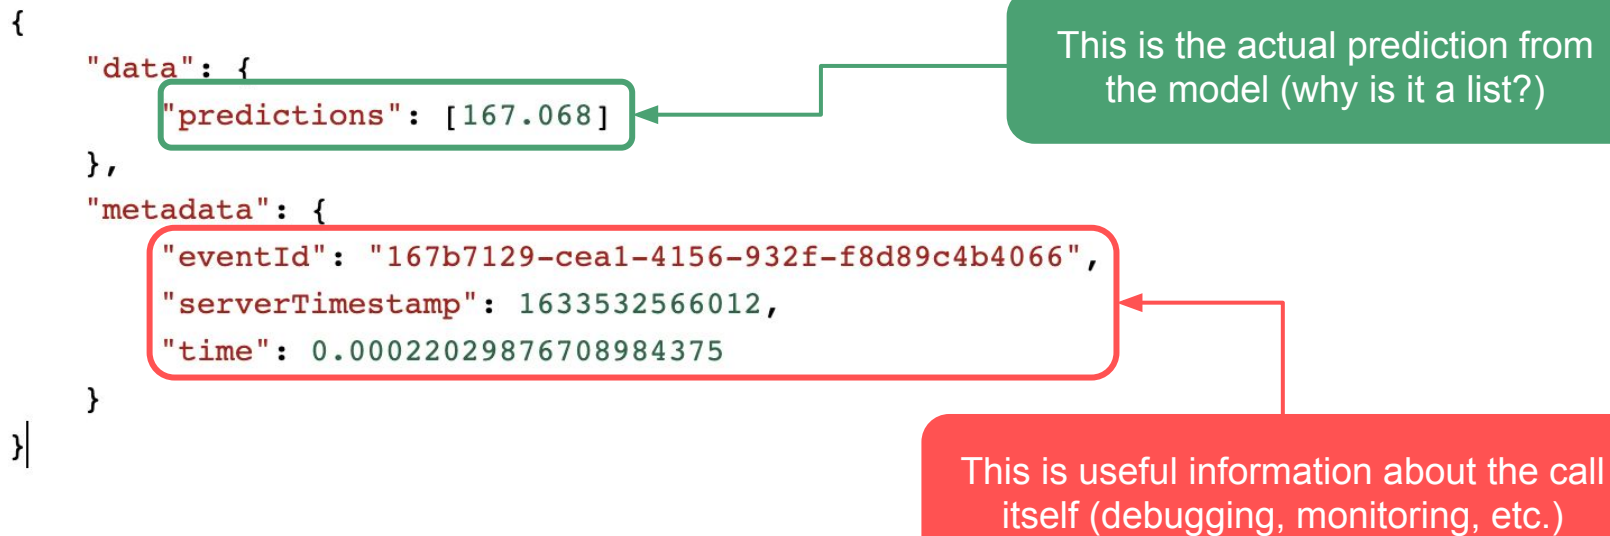
User asks for prediction

The full workflow

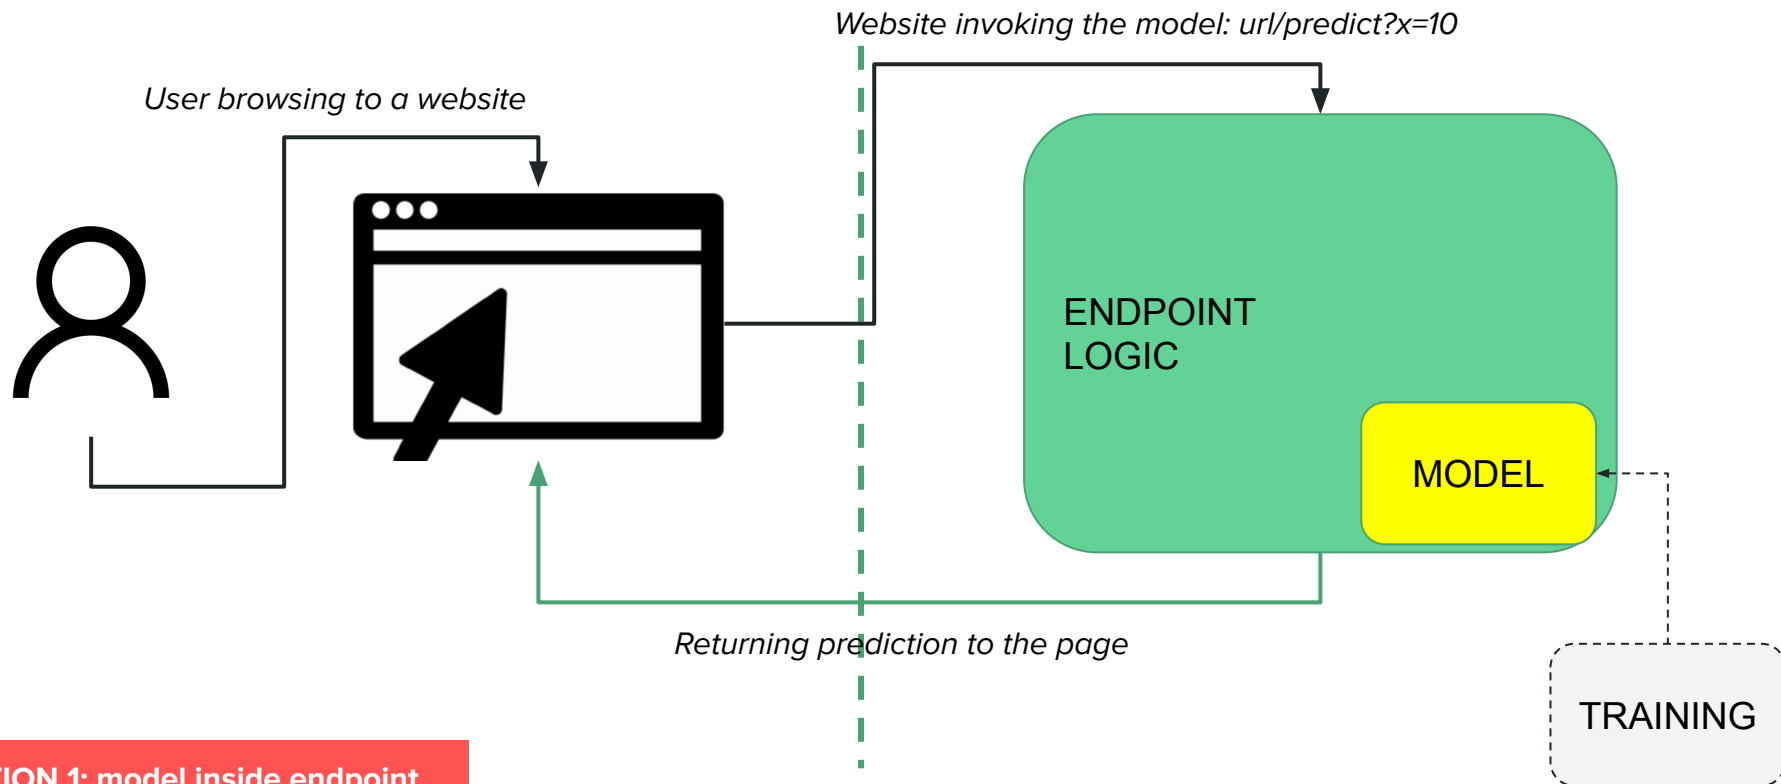


BONUS: Structuring the response

- Now **everybody** with the URL can use your awesome model!
- Can we make the response a bit clearer?

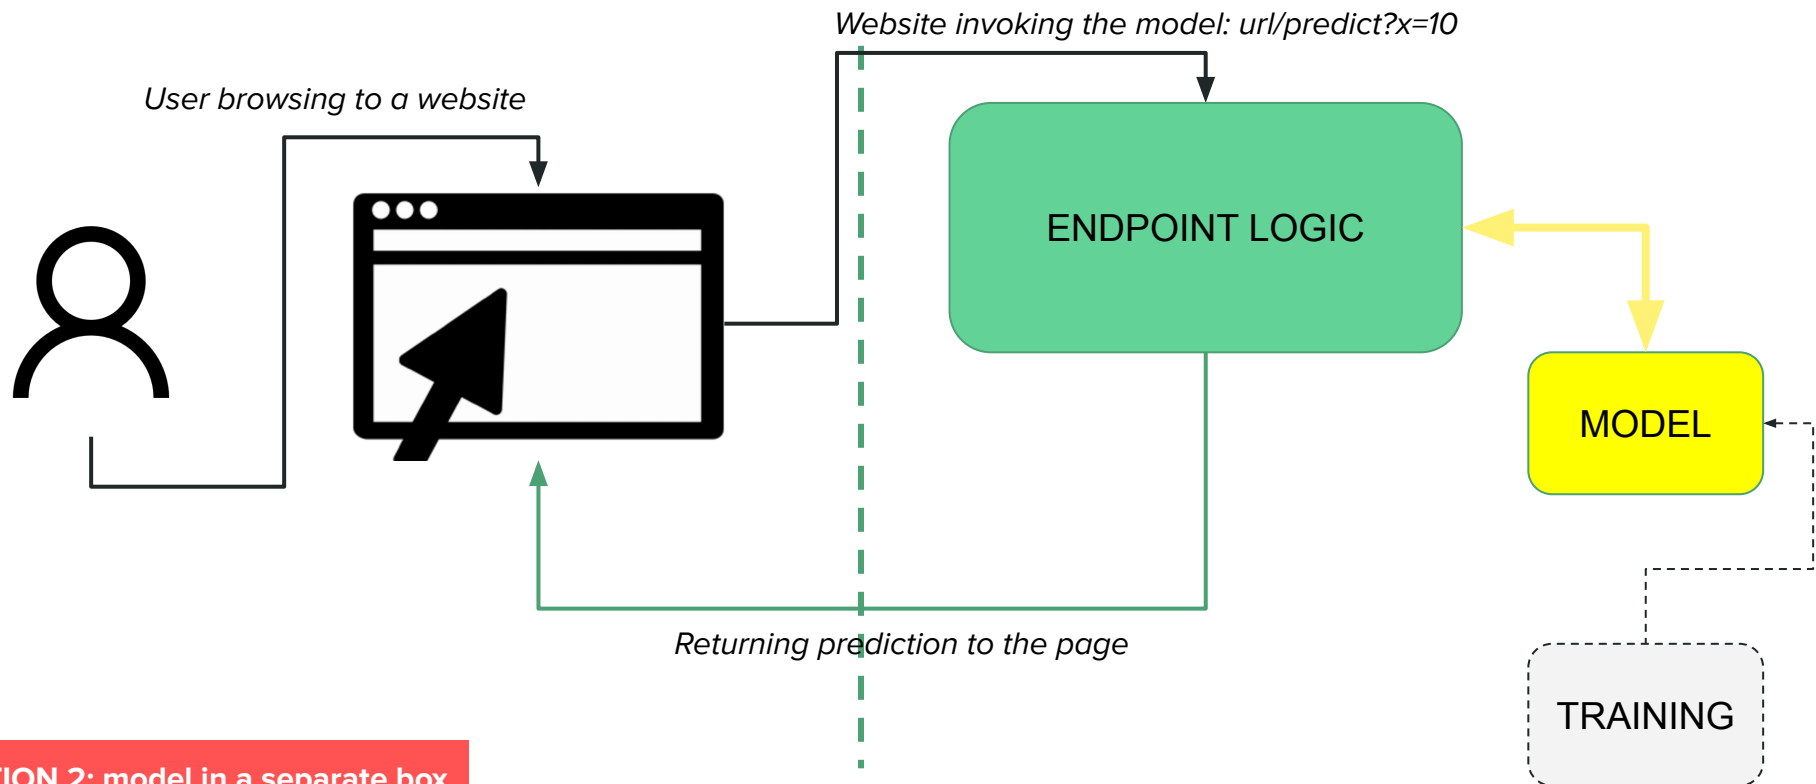


Scenario 1: Endpoint with Model (ours)



OPTION 1: model inside endpoint

Scenario 2: Endpoint + Model



OPTION 2: model in a separate box

Flask... in the cloud

Deploying ML models To the Web with Flask on AWS EC2 Instance



- There's 1M and one tutorials on how to use the very same tools (a Flask web app, a simple HTML + Javascript page) to port your app into the cloud.
- BONUS points for your final demo if you can show your endpoint in the cloud, either through an EC2 or Streamlit Cloud (**make sure to use the free resources / plan to avoid incurring in costs!**)

Alternative deployment scenarios

There is a ton of alternatives when it comes to *serving predictions* from the cloud, ranging from pure infrastructure to fully managed services. For example:

- You can deploy your model manually on a virtual machine, by installing Flask and run through screen (like they do [here](#))
- You can deploy your model through a web app hosted by Elasticbeanstalk (like they do [here](#))
- You can deploy your model through a web app hosted by Fargate (like they do [here](#))
- You can deploy your model through Sagemaker, and expose it through a lambda (like [we did in the 2021 repository](#))

After deployment: monitoring

We are not going to discuss monitoring, as we are not launching new apps in this course (for now!). However, after our model is live we need to:

- monitor how the pipeline is doing:
 - How is the new data coming in?
 - Does the model need re-training?
 - Is my new model better than the old one?
- check what users are doing with it!

After deployment: monitoring

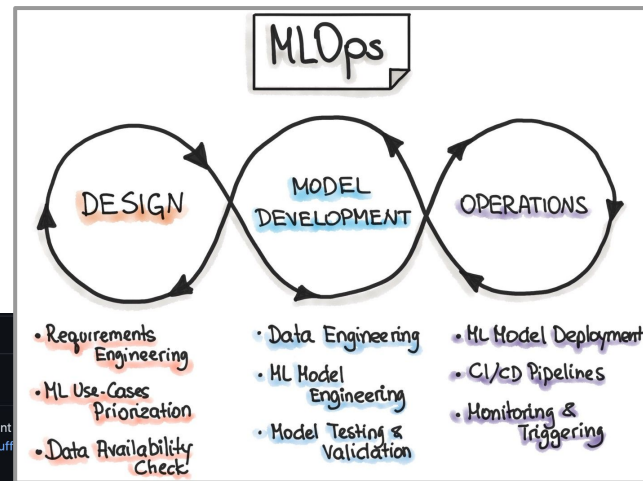
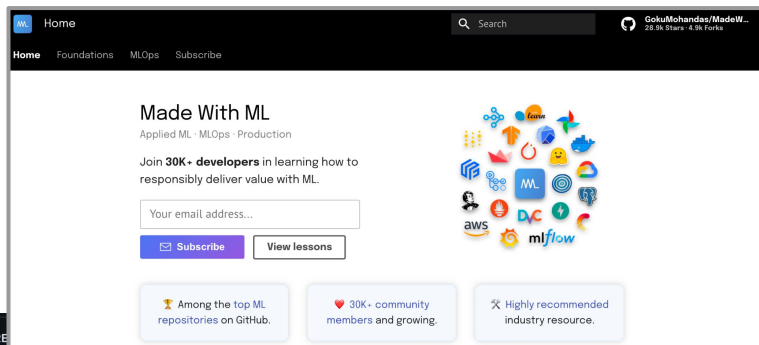
We are not going to discuss monitoring, as we are not launching new apps in this course (for now!). However, after our model is live we need to:

- monitor how the pipeline is doing:
 - How is the new data coming in?
 - Does the model need re-training?
 - Is my new model better than the old one?
- check what users are doing with it!
 - You never know how people would use stuff!



The adventure never stops!

There is a ton of recent developments in the “MLOps” space (we do our small part as well in the community). If you want to know more, reach out!



no-ops-machine-learning

A PaaS End-to-End ML Setup with Metaflow, Serverless and SageMaker.

This repo is an end-to-end model building exercise, showing how to go from a dataset file to an endpoint serving predictions through a repeatable, modular and scalable process (DAG-based). More importantly, all of this is achieved in minutes from a developer laptop, without explicitly deploying/maintaining *any* infrastructure.

For the full context and details about the tutorial, please refer to the [blog post](#).

Overview

README.md

You Don't Need a Bigger Boat

An end-to-end (Metaflow-based) implementation of an intent...

This is a WIP - check back often for updates.

Philosophical Motivations

There is plenty of tutorials and blog posts around the Internet on data pipelines and tooling. However:

- they (for good pedagogical reasons) tend to focus on *one* tool / step at a time, leaving us to wonder how the rest of the pipeline works;
- they (for good pedagogical reasons) tend to work in a toy-world fashion, leaving us to wonder what would happen when a real dataset and a real-world problem enter the scene.