

# NYU FRE 7773 - Week 12

---

*Machine Learning in Financial Engineering*

Jacopo Tagliabue


# RecSys 101 (again!)

---

*Machine Learning in Financial Engineering*  
Jacopo Tagliabue

# RecSys by use case

- RS can be understood easily by use case and input-output:



# RecSys by use case


- RS can be understood easily by use case and input-output:
  - Item as input, item as output: similar vs complementary

## Similar books based on genre




Sponsored ⓘ

Spring in Action, Sixth Edition  
Craig Walls  
★★★★★ 10  
Paperback  
\$53.99 ✓prime



Types and Programming Languages (The MIT Press)  
Benjamin C. Pierce  
★★★★★ 66  
Hardcover  
\$64.54 ✓prime



Essentials of Programming Languages, third edition (The MIT Press)  
Daniel P. Friedman  
★★★★★ 14  
Hardcover  
\$73.95 ✓prime

## Input

**Formal Semantics of Programming Languages**  
by Glynn Winskel (Author)  
★★★★☆ 17 ratings [Look inside ↴](#)

Hardcover	Paperback	Other Se
\$45.07	\$29.98 - \$70.00	See all 5 version
<input type="radio"/> Buy used: \$29.98		
<input checked="" type="radio"/> Buy new: \$70.00		
In Stock.		
Ships from and sold by Amazon.com.		
May be available at a lower price from other sellers, per Prime shipping.		


# RecSys by use case

- RS can be understood easily by use case and input-output:
  - Item as input, item as output: similar vs complementary

## Frequently bought together



+



+



Total price: **\$183.47**

[Add all three to Cart](#)

**i** Some of these items ship sooner than the others. [Show details](#)

- ✓ **This item:** Formal Semantics of Programming Languages by Glynn Winskel Paperback **\$70.00**
- ✓ Types and Programming Languages (The MIT Press) by Benjamin C. Pierce Hardcover **\$64.54**
- ✓ Practical Foundations for Programming Languages by Robert Harper Hardcover **\$48.93**

Input

**Formal Semantics of Programming Languages**  
by Glynn Winskel (Author)  
★★★★☆ 17 ratings [Look inside](#)

**Hardcover** \$45.07

**Paperback** \$29.98 - \$70.00

**Other Se** See all 5 versions

☐ Buy used: \$29.98

☒ Buy new: \$70.00

In Stock.


Ships from and sold by Amazon.com.

May be available at a lower price from other sellers, plus Prime shipping.

# RecSys by use case

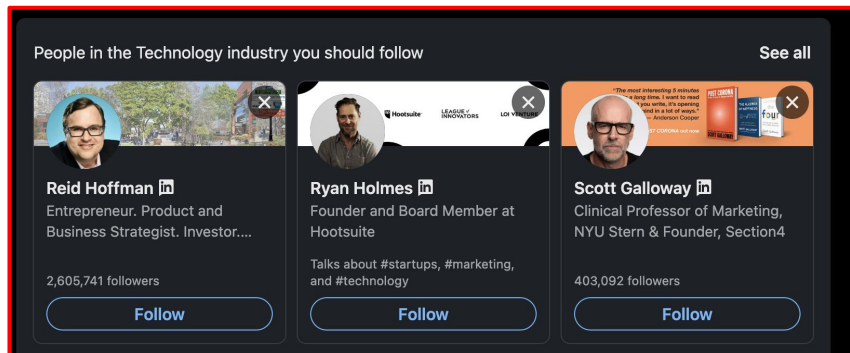
- RS can be understood easily by use case and input-output:
  - Item as input, item as output: similar vs complementary
  - User as input, item as output: “for you”

Input



# RecSys by use case

- RS can be understood easily by use case and input-output:
  - Item as input, item as output: similar vs complementary
  - User as input, item as output: “for you”
  - User as input, user as output: people you may know

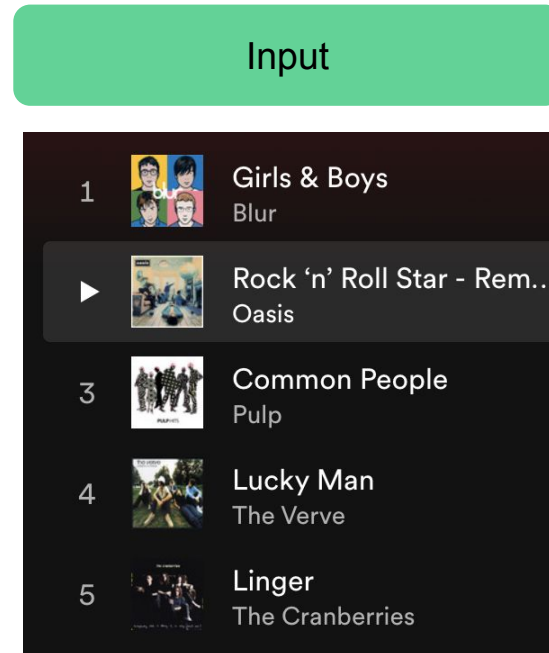



Input



# RecSys by use case

- RS can be understood easily by use case and input-output:
  - Item as input, item as output: similar vs complementary
  - User as input, item as output: “for you”
  - User as input, user as output: people you may know
  - Session as input, item as output: what are you doing next?





# RecSys by use case

- RS can be understood easily by use case and input-output:
  - Item as input, item as output: similar vs complementary
  - User as input, item as output: “for you”
  - User as input, user as output: people you may know
  - Session as input, item as output: what are you doing next?
  - Item as input, user as output: who should we sell this to?



Input

*New Fantastic SaaS  
Product!*

# RecSys by use case (with refs!)

- RS can be understood easily by use case and input-output:
  - Item as input, item as output: similar vs complementary
  - User as input, item as output: “for you”
  - User as input, user as output: people you may know
  - Session as input, item as output: what are you doing next?
  - Item as input, user as output: who should we sell this to?

*Key intuition:* if you like **X**, you like things similar to **X** as well!


What does “similar”  
mean?

# Similarity and representation

---

# Intuition: similarity is “closeness” in a “proper” space

- Let's map movies along two dimensions:
  - How much action is there?
  - How much violence is there?
- Observation #1: items as vectors
  - Indiana Jones: [ 3, 1 ]
  - Fast and Furious: [ 5, 2 ]
  - Die Hard: [ 4, 4 ]
- Observation #2: similar movies are close in the space
  - Back to RecSys: if you like Indiana Jones, you're more likely to like FF than Die Hard



# Intuition: similarity is “closeness” in a “proper” space

- If the space does not *represent* the underlying concepts well, we are in trouble!
- Machines understand vectors, but not all vectorizations define an appropriate space in this sense.
- For example, let’s consider one-hot encoding:
  - Is “cat” more similar to “dog” than “snake”?

Dog

1	0	0
---	---	---

Snake

0	1	0
---	---	---

Cat

0	0	1
---	---	---

## Intuition: similarity is “**closeness**” in a “proper” space

- If the space *represents* the underlying concepts well, items close in the space will be similar, items far apart are not so similar.
- While there are many **different ways** to characterize “close”, cosine distance (or dot product on scaled vectors) is the most common.
- **Corollary:** “similarity inference” is “just” nearest neighbor search in the vector space

$$\text{cosine}(\mathbf{v}, \mathbf{w}) = \frac{\mathbf{v} \cdot \mathbf{w}}{|\mathbf{v}| |\mathbf{w}|} = \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$


Dot Product



# Intuition: a (basic) recSys is like a GPS navigator

Consider a movie recommendation systems (user-item case)


- **Step 1:** represent movies in a suitable space



# Intuition: a (basic) recSys is like a GPS navigator

Consider a movie recommendation systems (user-item case)

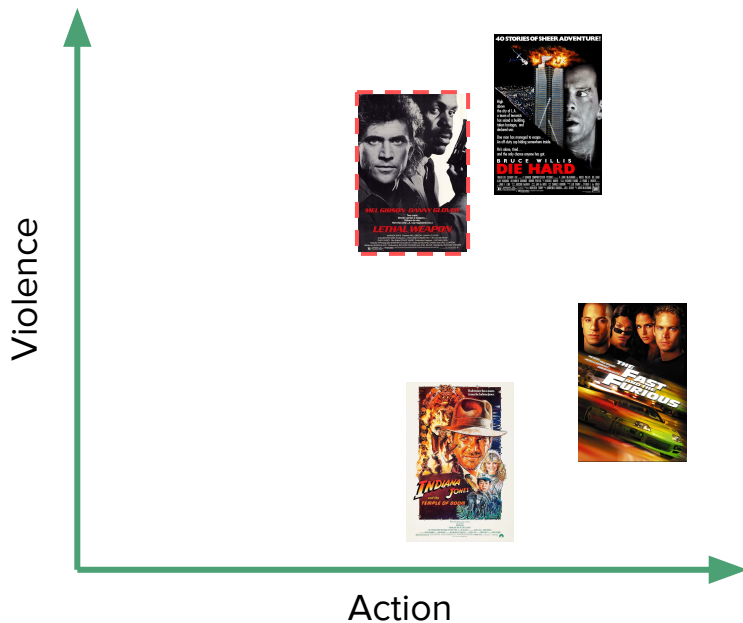
- **Step 1:** represent movies in a suitable space
- **Step 2:** represent the user *position* in the space
  - For example, Jacopo is “where his last movie is”



# Intuition: a (basic) recSys is like a GPS navigator

Consider a movie recommendation systems (user-item case)

- **Step 1:** represent movies in a suitable space
- **Step 2:** represent the user *position* in the space
- **Step 3:** recommend the closest K items (KNN search) to the user!
  - Recommendation: Die Hard!




# Intuition: a (basic) recSys is like a GPS navigator

Consider a real life example:

- **Jacopo** goes on vacation in Maui, Hawaii
  - Does Jacopo like surfing?
- **Ethan** goes on vacation in Boulder, Colorado
  - Does Ethan like climbing?
- **Intuition:** by knowing the position of the users in the space (in this case, Earth), we can tell a lot about their preferences!

**A huge part in our success when building recSys boils down to the quality of the representation in our space.**

***How do we map users and items to vectors, then?***



# Representations and data sources

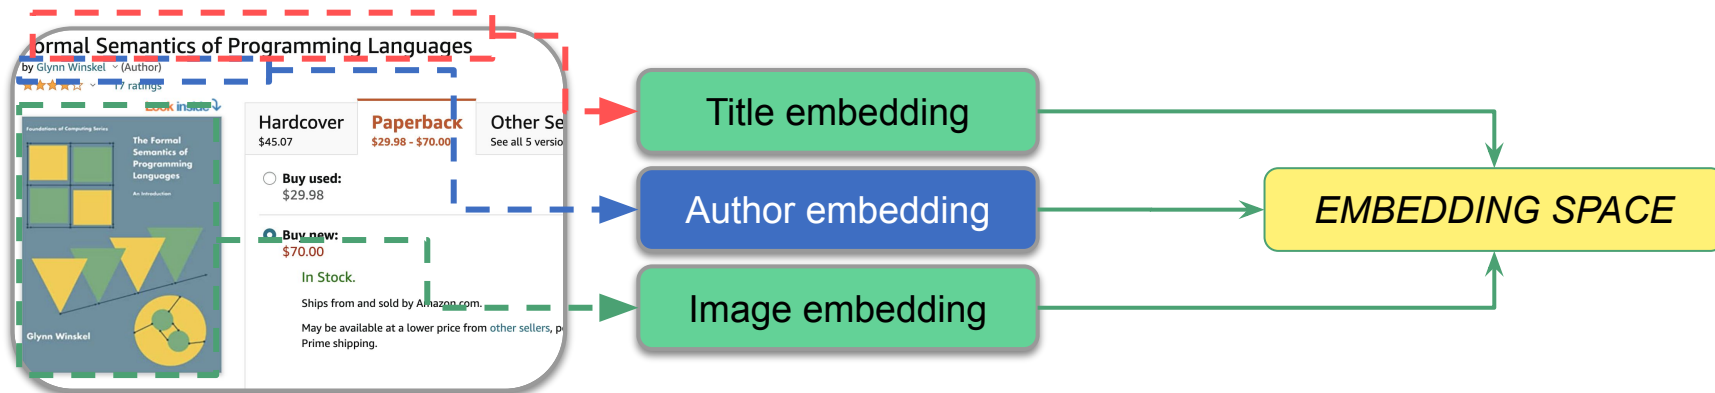
- **GOAL:** learning a good representation space!
  - A “good” space is a space where items that are indeed similar are close, and items that are far from each other are unrelated.
- While of course we could ask humans to rate *action vs violence vs comedy* ... for all movies on Netflix, that is impractical:
  - A ton of manual work (imagine doing this for all the books on Amazon!)
  - Unclear where to stop: should we have a dimension for actors as well? What about movie length? What about cost of production? Etc.
- We typically distinguish between **content-based** and **behavioral-based** representations (of course, hybrid are also possible!): e.g. for Netflix
  - Content: analyze the title, script, images from the movie, genre etc. - i.e. *what do we know about this item in our catalog?*
  - Behavioral: analyze the behavior of users wrt the items - **if users 1 like items A and B, and then likes also C, can we suppose C is similar to A and B?**

# Representations and data sources

- **Content-based** representations require only the “catalog” of the target entities in our RecSys scenario: their quality depends typically on their ability to turn images, text and categories into “good” vectors.
- **Behavioral-based** representations require real-world data from “users”, e.g.:
  - Purchase data from Amazon
  - Streaming data from Netflix
  - Playlist data from Spotify
  - etc.
- **Note:** a huge lesson of the last 20 years in RecSys is that behavioral-based representations are surprisingly useful in producing good representations (i.e. there is a lot of signals in people behavior!).
  - Q: when a behavioral strategy won't be helpful (we discussed it in class!)?

# Representations and data sources

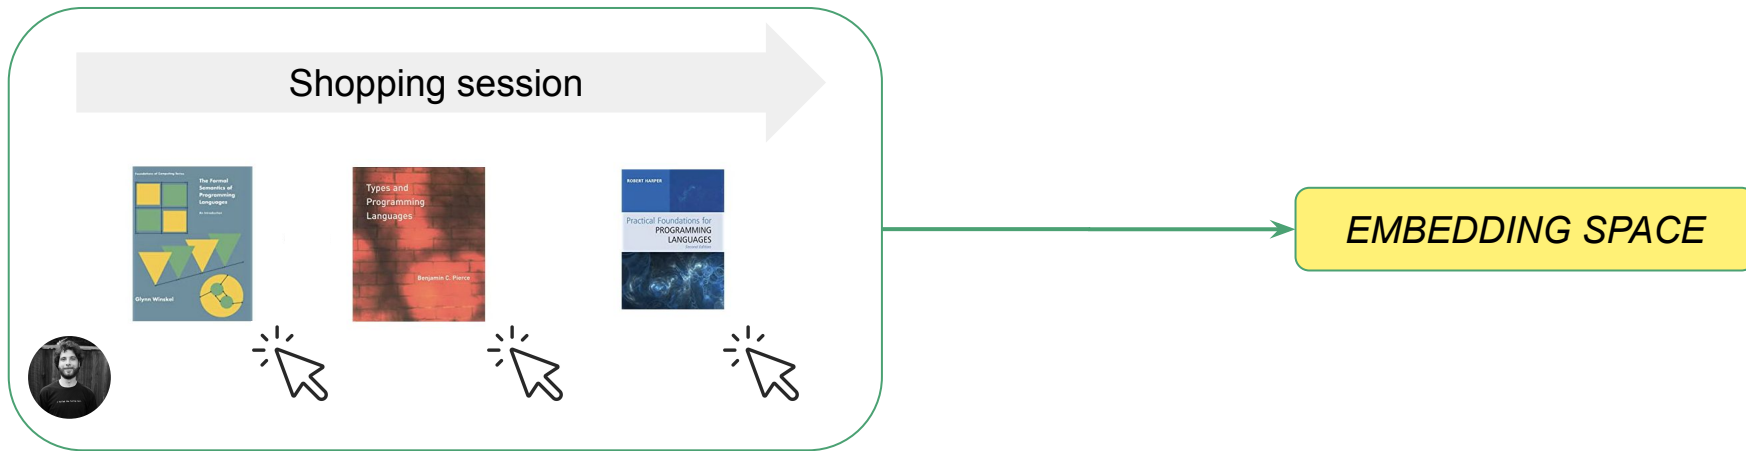
- We call the process of mapping high-cardinality entities (users and items, but also words etc.) to a low-dimensional space “**embedding**”.



*Content-based Embedding Space*

# Representations and data sources

- We call the process of mapping high-cardinality entities (users and items, but also words etc.) to a low-dimensional space “**embedding**”.



*Behavioral-based Embedding Space*



Word2Vec, Song2Vec, Everything2Vec

---

# The NLP Analogy: similar things appear often together

- **Distributional hypothesis:** “words that appear in similar contexts have similar meanings”
  - Example 1: if two books are often viewed in the same shopping session, they are probably similar!
  - Example 2: if two songs are often after each other in playlists, they are probably similar!

# The NLP Analogy: similar things appear often together

- **Distributional hypothesis:** “words that appear in similar contexts have similar meanings”
- Computational hypothesis: learn a classifier that given a target word (e.g. cat) tell me how likely is that a context word appear next to it (**Germany**, **furry**), *and use the learned weights as the word vector*. Since weights adjust during training to make sure similar words have high probability, their vectors will be close in the resulting embedding space.

**Word Embeddings**  
Past, Present and Future


# A recipe for learning word embeddings (“word2vec”)

- Outline of the general argument:
  - We need to learn vectors for words to make a “good” space
  - Words which are similar tend to appear in the same sentence
  - If we use vectors as weights for a classifier that tells when two words are likely to appear together, we can learn vectors that encode similarity and therefore produce a good space!
- In other words, the distributional hypothesis gives us a proxy measure of similarity: embeddings that are good in the distributional settings SHOULD therefore be good representations for word similarity.

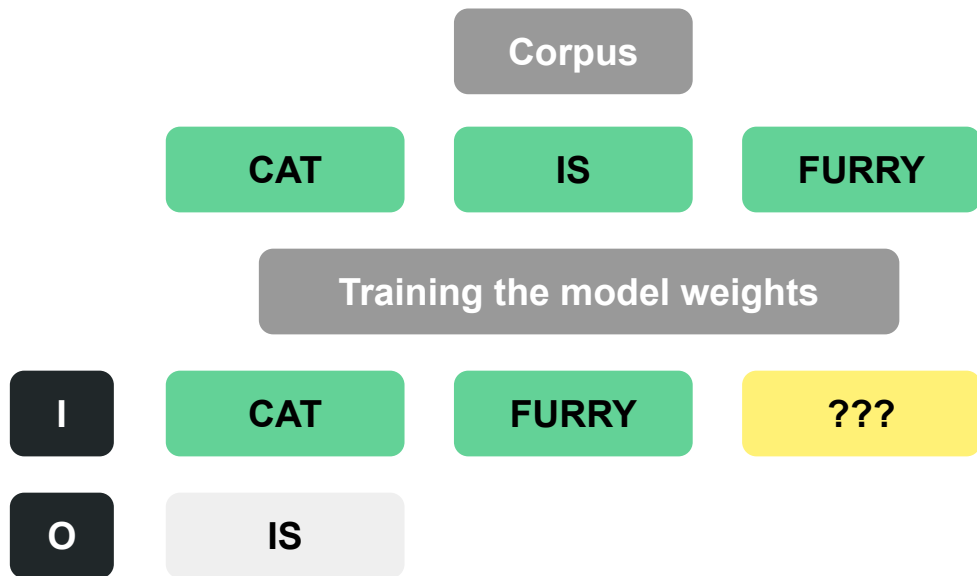
# A recipe for learning word embeddings (“word2vec”)

- Outline of the general argument:
  - We need to learn vectors for words to make a “good” space
  - Words which are similar tend to appear in the same sentence
  - If we use vectors as weights for a classifier that tells when two words are likely to appear together, we can learn vectors that encode similarity and therefore produce a good space!
- In other words, the distributional hypothesis gives us a proxy measure of similarity: embeddings that are good in the distributional settings SHOULD therefore be good representations for word similarity.

**BONUS POINT:** *this actually works in most successful applications of deep learning - if you find an appropriate prediction task, you can learn good representations for anything!*

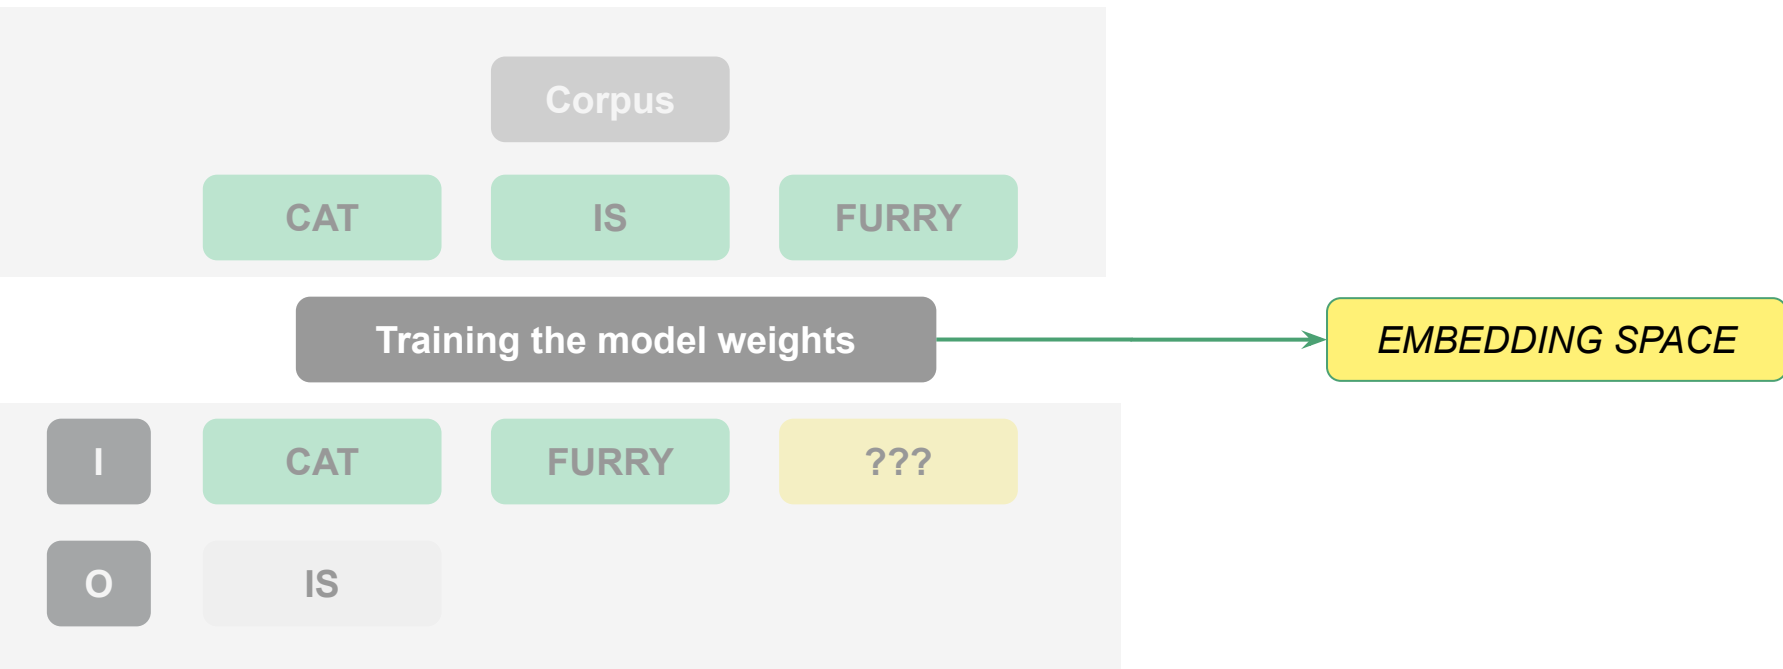


# A recipe for learning word embeddings (“word2vec”)



*Turn sequential data into a prediction problem*

# A recipe for learning word embeddings (“word2vec”)



*Weights associated with words are our embeddings*

# Word vectors in a *prediction* task

- CORPUS: “The furry cat is on the mat”
- WINDOW LENGTH: 2
- TARGET: “cat”
- INPUT PREPARATION, positive and negative samples

Target	Context	Label
cat	furry	1
cat	the	1
cat	is	1
cat	on	1

Target	Context	Label
cat	Berlin	0
cat	Jacopo	0
cat	ciao	0
cat	table	0



# Word vectors in a *prediction* task

- CORPUS: “The furry cat is on the mat”
- WINDOW LENGTH: 2
- TARGET: “cat”
- INPUT PREPARATION, positive and negative samples (a=0.75)

$$P_{\alpha}(w) = \frac{\text{count}(w)^{\alpha}}{\sum_{w'} \text{count}(w')^{\alpha}}$$

Target	Context	Label
cat	furry	1
cat	the	1
cat	is	1
cat	on	1

Target	Context	Label
cat	Berlin	0
cat	Jacopo	0
cat	ciao	0
cat	table	0

# Word vectors in a *prediction* task

- We have turned a word prediction problem into a binary classification problem
  - Is the context word likely to appear next to the target word?
- Let's define our learning objective:
  - We want to maximize the similarity of  $(t,c)$  drawn from the positive examples
  - We want to minimize the similarity of  $(t,c)$  drawn from the negative examples

$$L(\theta) = \sum_{(t,c) \in +} \log P(+|t,c) + \sum_{(t,c) \in -} \log P(-|t,c)$$

CHAPTER

6

## Vector Semantics and Embeddings

荃者所以在鱼，得鱼而忘荃 Nets are for fish;  
Once you get the fish, you can forget the net.  
言者所以在意，得意而忘言 Words are for meaning;  
Once you get the meaning, you can forget the words  
庄子(Zhuangzi), Chapter 26

The parable that Lao Anshu is famous for comes mainly on its forearm. But

# Word vectors in a *prediction* task

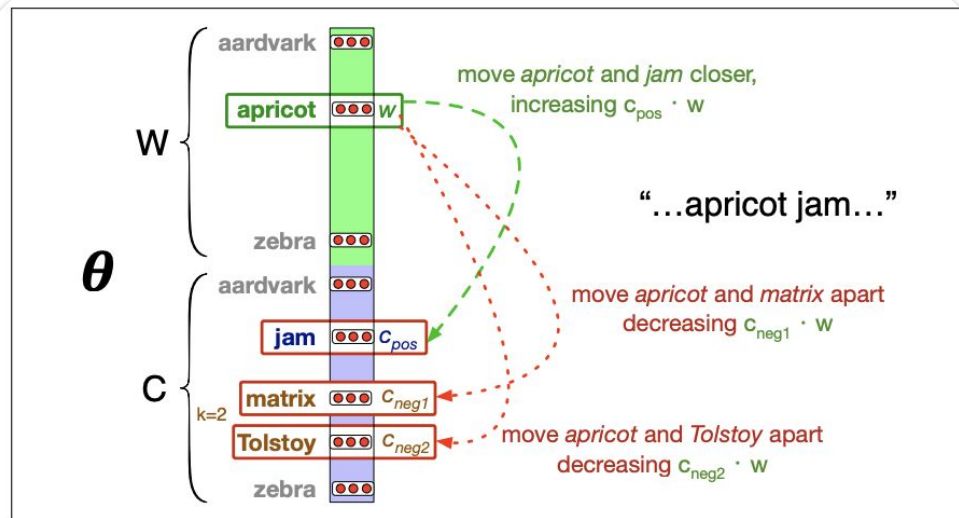
- Let's define our learning objective:
  - We want to maximize the similarity of (t,c) drawn from the positive examples
  - We want to minimize the similarity of (t,c) drawn from the negative examples

The diagram illustrates the components of the loss function. A blue box labeled "Dot product" has an arrow pointing to the term  $c \cdot t$  in the first equation. A green box labeled "Sigmoid" has an arrow pointing to the  $\sigma$  function in the same term. The equations are as follows:

$$= \log \sigma(c \cdot t) + \sum_{i=1}^k \log \sigma(-n_i \cdot t)$$
$$= \log \frac{1}{1 + e^{-c \cdot t}} + \sum_{i=1}^k \log \frac{1}{1 + e^{n_i \cdot t}}$$

# Word vectors in a *prediction* task


- **Remember:** we maximize the dot product of the word with the context words, and minimize the dot products of the word with the negative sampled words!
- Training procedure:
  - Random initialization of vectors (embeddings) for N words in the vocabulary.
  - At each step, move embeddings of related words closer in the vector space, and push others further away (using gradient descent).



**Figure 6.14** Intuition of one step of gradient descent. The skip-gram model tries to shift embeddings so the target embeddings (here for *apricot*) are closer to (have a higher dot product with) context embeddings for nearby words (here *jam*) and further from (lower dot product with) context embeddings for noise words that don't occur nearby (here *Tolstoy* and *matrix*).


# Is this a “good” space?

- word2vec tends to capture well similarity between words and some analogical relations - **without any human labels / intervention!**
- Once you have a well-trained embedding space, the offsets between vector embeddings can be used to solve analogies such as: “man : king = women : ?” (*queen*)
  - This is possible since the result of  $\text{vector}(\text{'king'}) - \text{vector}(\text{'man'}) + \text{vector}(\text{'woman'})$  is a vector close to  $\text{vector}(\text{'queen'})$ .



# From NLP, back to RecSys

**Remember:** the same intuition about “words in a sentence” can be applied whenever we have meaningful sequences of target items (e.g. playlist, shopping sessions etc.)



CAT

IS

FURRY

*Song2Vec*



CAT

IS

FURRY

*Book2Vec*



Coding time!