

NYU FRE 7773 - Week 13

Machine Learning in Financial Engineering

Jacopo Tagliabue

Evaluation time
(sorry, need to do it)!

RecSys part II


Machine Learning in Financial Engineering
Jacopo Tagliabue

Recap on embeddings

Intuition: a (basic) recSys is like a GPS navigator

Consider a movie recommendation systems (user-item case)


- **Step 1:** represent movies in a suitable space



Intuition: a (basic) recSys is like a GPS navigator

Consider a movie recommendation systems (user-item case)

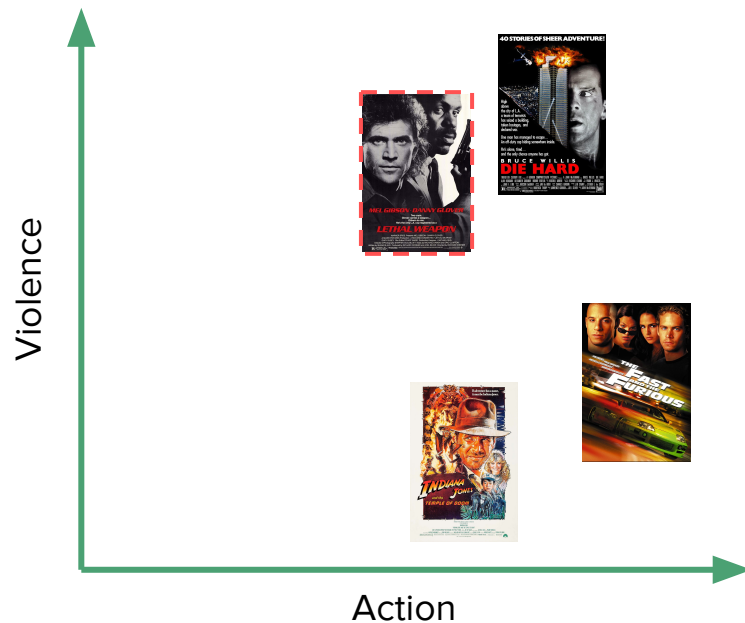
- **Step 1:** represent movies in a suitable space
- **Step 2:** represent the user *position* in the space
 - For example, Jacopo is “where his last movie is”



Intuition: a (basic) recSys is like a GPS navigator


Consider a movie recommendation systems (user-item case)

- **Step 1:** represent movies in a suitable space
- **Step 2:** represent the user *position* in the space
- **Step 3:** recommend the closest K items (KNN search) to the user!
 - Recommendation: Die Hard!




Word2Vec, Song2Vec, Everything2Vec


- **Word2Vec** builds embeddings for words, leveraging their vicinity in a sequence.
- The same intuition can be applied whenever we have meaningful sequences of target items (e.g. playlist, shopping sessions etc.)



CAT



IS



FURRY

Song2Vec



CAT

IS

FURRY




Book2Vec

Building a song recommender


RecSys Data, revisited

- Input (“features”): as we discussed we divide them typically in **content-based** vs **behavioral-based**.
- Target: typically two types, ratings (stars on a movie) and interactions (click on ads, buying shoes):
 - Reviews are *explicit* (people tell you when you suck!), interaction is *implicit* (people buy stuff they like, but don’t tell you what they don’t!). The target will influence your loss function, typically mean squared error for ratings, logistic loss for binary.
 - As many (all?) hard ML problems, recSys are fundamentally a challenge in sparsity: in the Netflix challenge for example, the users / movies matrix has **100 million / ~8.5 billion ratings**.



Once upon a time: collaborative filtering

- **Use case:** user-item, Netflix prize
- **Intuition:** we can “reduce” the original matrix (sparse) to the combination of two smaller matrices, users/movies
 - Instead of listing all my ratings, we could say I like sci-fi and Tom Cruise movies
- **Prediction:** the dot product between user and item vectors is the estimate of the rating!
- **Training:**
$$\min_{p,q} \sum_{(u,i) \in \kappa} (r_{ui} - q_i^T p_u)^2 + \lambda(\|q_i\|^2 + \|p_u\|^2)$$



Decomposing the matrix into user and movie features

Q: wait a sec, what about cold start?

Once upon a time: collaborative filtering

- **Collaborative filtering** is great, but its performance crucially depends on behavioral data. Cold-start scenarios are not easily addressed in the original problem formulation!
- While you can think of fixing the problem in clever ways (for example, you can ask “few good users” to help you out by explicitly asking them initial ratings), the underlying “bug” of pure matrix factorization is the lack of content-based features. If the only type of data the model is using is interactions / ratings / purchases, we are missing the chance of leveraging important meta-data about our use case!
- While there are pure matrix-based methods that can help, more modern systems leverage deep learning and its inherent flexibility.

Song2Vec

We build a Metaflow flow to recommend new songs: for example, *you are listening to “Hey Jude” and want to know what to listen next, or you’re building a playlist and want to know what to add.*

We need:

- Initial dataset
- Modelling assumption
- Train and test preparation
- The evaluation

The dataset

- Initial dataset: the Playlist Dataset from Spotify
 - Each playlist has a name and an array of tracks.
 - Each track contains basic metadata and a track id.

Introducing The Million Playlist Dataset and RecSys Challenge 2018



May 30, 2018
Published by Ching-Wei Chen

RecSys Challenge 2018

Welcome ACM RecSys Community! For this year's challenge, use the Spotify Million Playlist Dataset to help users create and extend their own playlists.

```
{
  "name": "musical",
  "collaborative": false,
  "pid": 5,
  "modified_at": 1493424000,
  "num_albums": 7,
  "num_tracks": 12,
  "num_followers": 1,
  "num_edits": 2,
  "duration_ms": 2657366,
  "tracks": [
    {
      "pos": 0,
      "artist_name": "Degiheugi",
      "track_uri": "spotify:track:7vqa3sDmtEaVJ2gcvxtRID",
      "artist_uri":
"spotify:artist:3V2paBXEoZIAhfZRJmo2jL",
      "track_name": "Finalement",
      "album_uri": "spotify:album:2KrRMJ9z7Xjoz1Az4O6UML",
      "duration_ms": 166264,
      "album_name": "Dancing Chords and Fireflies"
    },
    {
      "pos": 1,
      "artist_name": "Degiheugi"
```

Modelling assumptions

- **Word2Vec** = “**words** that appear together in **sentences** are similar”.
- **Song2Vec** = “**songs** that appear together in **playlists** are similar”.
- *We can apply Word2Vec to songs!*

ORDERS MATTER!

CAT

IS

FURRY

1



Girls & Boys
Blur



Rock 'n' Roll Star - Rem...
Oasis

3




Common People
Pulp

Modelling assumptions

- **Recommendation / Prediction**

- **Input:** a list (possibly $n=1$) of songs in a playlist / user history
- **Output:** predict the next song to listen / continue the playlist



Data preparation

- We divide track data from metadata instead of replicating track metadata every time a track is found:
 - Playlists -> sequences of track ids -> ["fasfa", "fauofu", "earfao"]
 - Catalog -> a mapping between track ids and metadata ->
 - "fauofu": { "name": "Hey Jude" }
 - "earfao": { "name": "Imagine" }

```
"name": "musical",
"collaborative": "false",
"pid": 5,
"modified_at": 1493424000,
"num_albums": 7,
"num_tracks": 12,
"num_followers": 1,
"num_edits": 2,
"duration_ms": 2657366,
"tracks": [
  {
    "pos": 0,
    "artist_name": "Degiheugi",
    "track_uri": "spotify:track:7vqa3sDmtEaVJ2gcvxtRID",
    "artist_uri":
"spotify:artist:3V2paBXEoZIAhfZRJmo2jL",
    "track_name": "Finalement",
    "album_uri": "spotify:album:2KrRMJ9z7Xjoz1Az4O6UML",
    "duration_ms": 166264,
    "album_name": "Dancing Chords and Fireflies"
  },
  {
    "pos": 1,
    "artist_name": "Degiheugi"
```

Data preparation

- Instead of doing it manually, we rely on the dataset prepared by RecList, an open source library for evaluating Recommender Systems.
 - Using RecList we get a training set and data_catalog ready to use!

```
def prepare_dataset(self):  
    """  
    Get the Spotify dataset in a convenient shape by using the  
    abstractions provided by RecList (https://reclist.io/)  
  
    NOTE: THIS MAY TAKE A WHILE THE FIRST TIME AS THE DATASET GETS DOWNLOADED!  
    """  
    from reclist.datasets import SpotifyDataset  
    # get the Spotify million playlist dataset as a RecDataset object  
    self.spotify_dataset = SpotifyDataset(force_download=False)  
    # check the dataset by printing out the first playlist  
    print("First playlist is", self.spotify_dataset.x_train[0][0]['playlist_name'])  
    self.data_catalog = self.spotify_dataset.catalog  
    test_track = self.data_catalog[list(self.data_catalog.keys())[0]]  
    print("First track metadata is", test_track['track_name'], test_track['artist_name'])
```

RecList

RecList is an open source library providing behavioral, "black-box" testing for recommender systems. Inspired by the pioneering work of Ribeiro et al. 2020 in NLP, we introduce a general plug-and-play procedure to scale up behavioral testing, with an easy-to-extend interface for custom use cases.

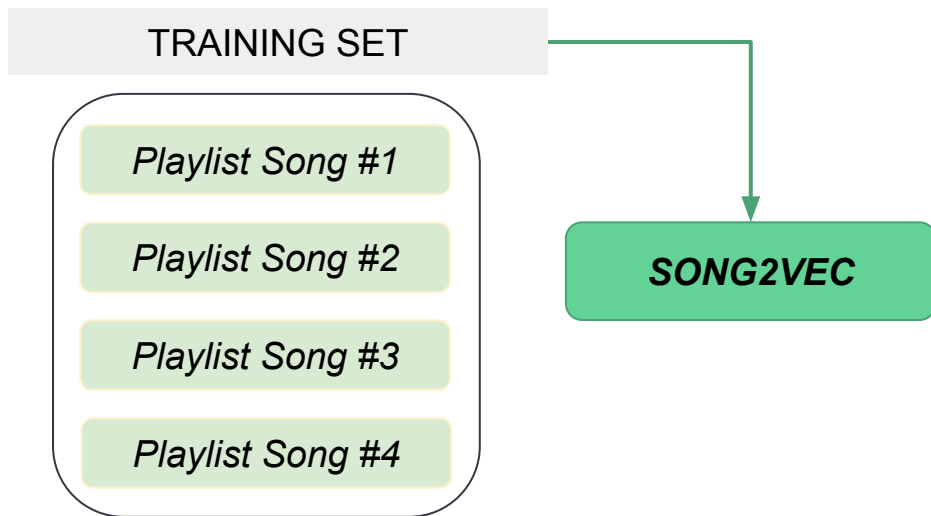
To streamline comparisons among existing models, RecList ships with popular datasets and ready-made behavioral tests: read the our [TDS blog post](#) as a gentle introduction to the main use cases, and try out our [colab](#) to get started with the code.

We are actively working towards our beta, with new

```
from reclist.datasets import CoveoDataset  
from reclist.recommenders.prod2vec import CoveoP2VRecModel  
from reclist.reclist import CoveoCartRecList  
  
coveo_dataset = CoveoDataset()  
  
model = CoveoP2VRecModel()  
model.train(coveo_dataset.x_train)  
  
# instantiate rec_list object  
rec_list = CoveoCartRecList(  
    model=model,  
    dataset=coveo_dataset
```

Train and test preparation

- We leverage RecList, which automatically treats the Spotify dataset as a sequential problem ready to be solved with a **Song2Vec** approach.
- We can prepare a training dataset *without labels*!



RecList 

RecList is an open source library providing behavioral, "black-box" testing for recommender systems. Inspired by the pioneering work of [Ribeiro et al. 2020](#) in NLP, we introduce a general plug-and-play procedure to scale up behavioral testing, with an easy-to-extend interface for custom use cases.

To streamline comparisons among existing models, RecList ships with popular datasets and ready-made behavioral tests: read the our [TDS blog post](#) as a

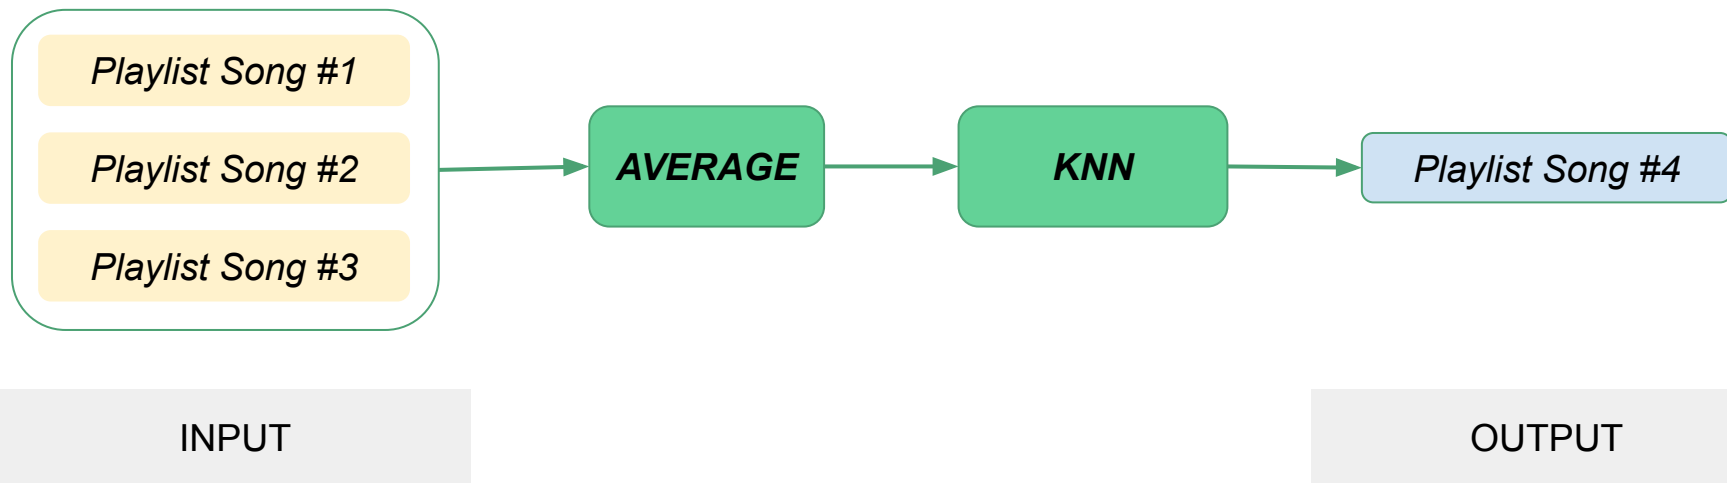
Train and test preparation

- *Input:* after training, we use the **average** of the embeddings in the first K songs as input for prediction.
- *Output:* we take the songs (through KNN) that are closer to that vector.

RecList 

RecList is an open source library providing behavioral, "black-box" testing for recommender systems. Inspired by the pioneering work of [Ribeiro et al. 2020](#) in NLP, we introduce a general plug-and-play procedure to scale up behavioral testing, with an easy-to-extend interface for custom use cases.

To streamline comparisons among existing models, RecList ships with popular datasets and ready-made behavioral tests: read the our [TDS blog post](#) as a



Train and test preparation


- Once predictions are made, we can run evaluation with the usual “word2vec” methodology, i.e. we compare our predictions with the target song (the last song in a playlist).

RecList 

RecList is an open source library providing behavioral, “black-box” testing for recommender systems. Inspired by the pioneering work of [Ribeiro et al. 2020](#) in NLP, we introduce a general plug-and-play procedure to scale up behavioral testing, with an easy-to-extend interface for custom use cases.

To streamline comparisons among existing models, RecList ships with popular datasets and ready-made behavioral tests: read the our [TDS blog post](#) as a

EVALUATION



Evaluation

- We can evaluate our recommender model with a combination of quantitative and behavioral tests, using RecList once again.

RecList

RecList is an open source library providing behavioral, "black-box" testing for recommender systems. Inspired by the pioneering work of [Ribeiro et al. 2020](#) in NLP, we introduce a general plug-and-play procedure to scale up behavioral testing, with an easy-to-extend interface for custom use cases.

To streamline comparisons among existing models, RecList ships with popular datasets and ready-made behavioral tests: read the our [TDS blog post](#) as a


```
from reclist.reclist import SpotifySessionRecList
# get the current value of the hyper in the foreach
self.window_length = int(self.input)
# instantiate and train the skip-gram model
model = SpotifyModel()
model.train(self.spotify_dataset.x_train)
print("Training with window {} is completed!".format(self.window_length))
# finally, version the embeddings as key values
self.track_vectors = model._model
rec_list = SpotifySessionRecList(
    model=model,
    dataset=self.spotify_dataset
)
rec_list(verbose=True)
```

Model

Dataset

Evaluation

- We can evaluate our recommender model with a combination of quantitative and behavioral tests, using RecList once again.
 - For example, we compute the distribution of hit-rate across various slices of data.



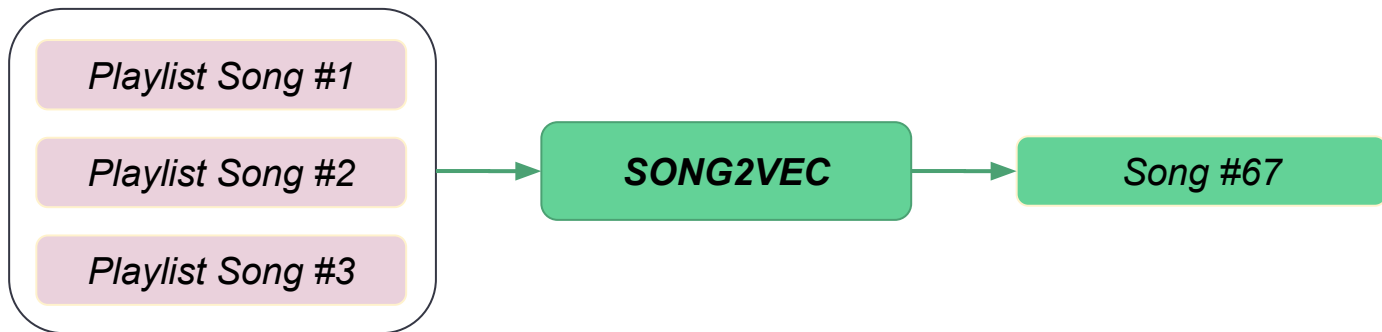
RecList

RecList is an open source library providing behavioral, "black-box" testing for recommender systems. Inspired by the pioneering work of [Ribeiro et al. 2020](#) in NLP, we introduce a general plug-and-play procedure to scale up behavioral testing, with an easy-to-extend interface for custom use cases.

To streamline comparisons among existing models, RecList ships with popular datasets and ready-made behavioral tests: read the our [TDS blog post](#) as a

Evaluation

- We can evaluate our recommender model with a combination of quantitative and behavioral tests, using RecList once again.
 - For example, we could check the system robustness to “perturbation”.



RecList

RecList is an open source library providing behavioral, “black-box” testing for recommender systems. Inspired by the pioneering work of [Ribeiro et al. 2020](#) in NLP, we introduce a general plug-and-play procedure to scale up behavioral testing, with an easy-to-extend interface for custom use cases.

To streamline comparisons among existing models, RecList ships with popular datasets and ready-made behavioral tests: read the our [TDS blog post](#) as a

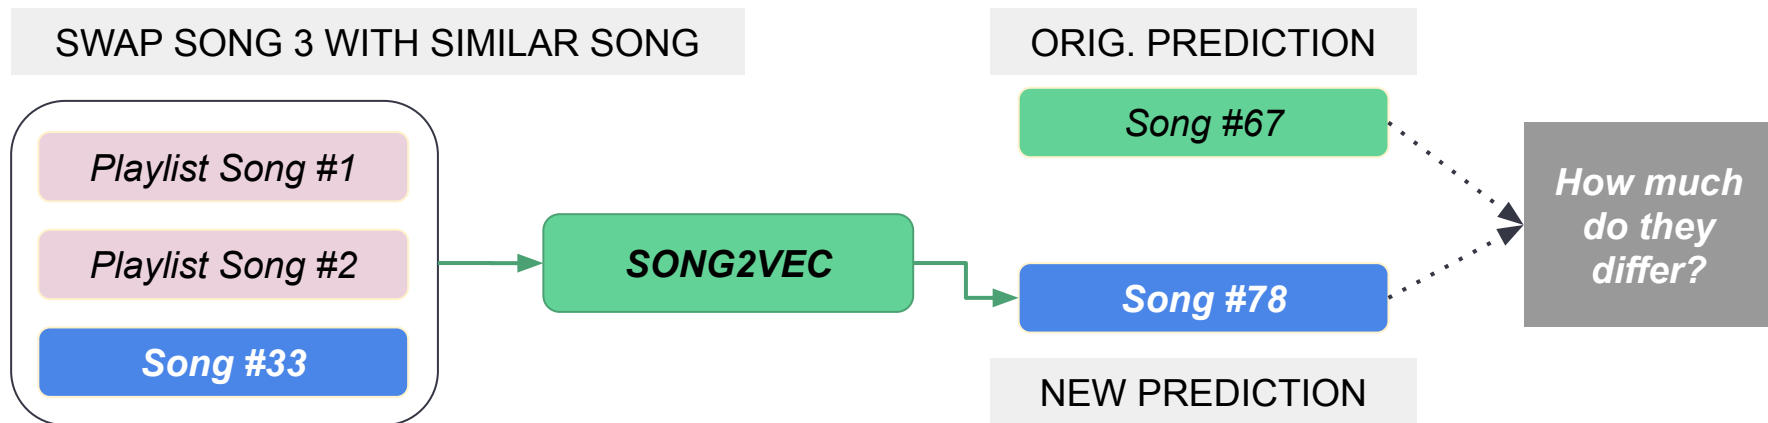
Evaluation

- We can evaluate our recommender model with a combination of quantitative and behavioral tests, using RecList once again.
 - For example, we could check the system robustness to “perturbation”.

RecList 

RecList is an open source library providing behavioral, “black-box” testing for recommender systems. Inspired by the pioneering work of [Ribeiro et al. 2020](#) in NLP, we introduce a general plug-and-play procedure to scale up behavioral testing, with an easy-to-extend interface for custom use cases.

To streamline comparisons among existing models, RecList ships with popular datasets and ready-made behavioral tests: read the our [TDS blog post](#) as a



What did Song2Vec learn?

Did we learn a “good” space?

- We evaluated our recommender models on typical RecSys metrics:
 - Hit rate, MRR, sliced-based tests, etc.
- By “solving” the prediction problem, we obtained embeddings for songs, which could be used *also* as features for other models, *provided that these embeddings are “good”*.
 - What can we do to check the quality of the representational space?

Fantastic Embeddings and How to Align Them: Zero-Shot Inference in a Multi-Shop Scenario

Federico Bianchi*
Bocconi University
Milano, Italy
f.bianchi@unibocconi.it

Jacopo Tagliabue*[†]
Coveo Labs
New York, NY
jtagliabue@coveo.com

Bingqing Yu*
Coveo
Montreal, Canada
cyu2@coveo.com

Luca Bigon*[‡]
Coveo
Montreal, Canada
lbigon@coveo.com

Ciro Greco*[§]
Coveo Labs
New York, NY
cgreco@coveo.com

ABSTRACT

This paper addresses the challenge of leveraging multiple embedding spaces for multi-shop personalization, proving that zero-shot inference is possible by transferring shopping intent from one website to another without manual intervention. We detail a machine learning pipeline to train and optimize embeddings *within shops* first, and support the quantitative findings with additional qualitative insights. We then turn to the harder task of using learned embeddings *across shops*: if products from different shops live in the same vector space, user intent - as represented by regions in this space - can then be transferred in a zero-shot fashion across websites. We propose and benchmark unsupervised and supervised methods to “travel” between embedding spaces, each with its own assumptions on data quantity and quality. We show that zero-shot personalization is indeed possible at scale by testing the shared embedding space with two downstream tasks, event prediction and two-shop suggestions. Finally, we create a space where

ACM Reference Format:


Federico Bianchi, Jacopo Tagliabue, Bingqing Yu, Luca Bigon, and Ciro Greco. 2020. Fantastic Embeddings and How to Align Them: Zero-Shot Inference in a Multi-Shop Scenario. In *Proceedings of ACM SIGIR Workshop on eCommerce (SIGIR eCom'20)*. ACM, New York, NY, USA, 11 pages.

1 INTRODUCTION

Inspired by the similarity between words in sentences and products in browsing sessions, recent work in recommender systems re-adapted the NLP CBOW model [20] to create *product embeddings* [17], i.e. low-dimensional representations which can be used alone or fed to downstream neural architectures for other machine learning tasks. Product embeddings have been mostly investigated as static entities so far, but, exactly as words [10], products are all but static. Since the creation of embeddings is a stochastic process, training embeddings for similar products in different digital shops


Did we learn a “good” space?

- Remember: a “good” space is a space where similar songs are close to each other!
- First check is a qualitative check: if we can label our entities (songs in this case) with human-readable concepts, we can inspect the space (with TSNE) to see if it matches our intuitions.
- **Example #1:** by tagging tracks with their artists, we can see beyonce and maroon 5 overlap more than beyonce and eminem.



Did we learn a “good” space?

- Remember: a “good” space is a space where similar songs are close to each other!
- First check is a qualitative check: if we can label our entities (songs in this case) with human-readable concepts, we can inspect the space (with TSNE) to see if it matches our intuitions.
- **Example #2:** by tagging tracks with keywords in playlists, *rock* and *pop* are clustered naturally, *wedding* songs overlap both.



Did we learn a “good” space?

- Second check is a quantitative check: if we can label our entities with independently established categories, we can train a classifier: input song vector, output category.
- **Example:** we can train a multi-label classifier that takes as input the embeddings we trained, and classify the artist for the song.
 - Intuitively, the better the classification, the more it “resembles” our intuitions.
 - The “accuracy number” *per se* is not meaningful, but this may be useful *to compare between different embedding spaces*.

Classification Report				
	precision	recall	f1-score	support
0	1.00	0.96	0.98	28
1	0.87	0.93	0.90	28
2	1.00	1.00	1.00	29
3	0.98	0.94	0.96	47
4	0.99	0.99	0.99	283
accuracy			0.98	415
macro avg	0.97	0.96	0.97	415
weighted avg	0.98	0.98	0.98	415

0 drake, 1 maroon 5, 2 eminem,
3 beyoncé, 4 unknown.

More embeddings!

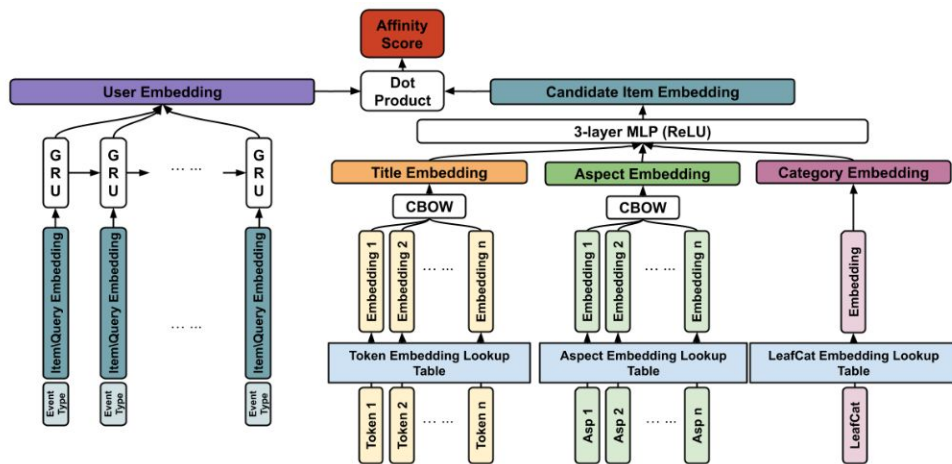
The two-tower model

- **Use case:** user-item
- **Intuition:** in Song2Vec, we represent the user *through the item embeddings* - but what if we can learn embeddings for user and items directly?
- **Prediction:** As usual with embeddings, we need to find a problem, such as, by predicting it, we will generate good vectors. In this case it is very direct, we use *existing interactions* data and try to predict unseen user-item interactions
- **Cold-start:** since neural nets are composable (i.e. backpropagation “just works” irrespective of input type), we can embed meta-data for users and items to help the model with cold start!

A two-tower model is composable!

The two-tower model

- A worked-out example from eBay
 - Users are represented through embeddings that consider previous items viewed and search queries
 - Products are represented with a combination of text and taxonomy information.



Coding time!